

A project report on

WEAPON DETECTION WITH EXPLAINABLE AI

Submitted partial fulfillment for the award of the degree of

Bachelor of Technology in Computer Science and Engineering

by

Aakarsh Ranjith 21BCE1205

Devika Rajeev 21BCE1193

Shahin Sha S 21BCE1885



VIT[®]

Vellore Institute of Technology
(Deemed to be University under section 3 of UGC Act, 1956)
CHENNAI

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

April, 2025

A project report on

WEAPON DETECTION WITH EXPLAINABLE AI

Submitted partial fulfillment for the award of the degree of

Bachelor of Technology in Computer Science and Engineering

by

Aakarsh Ranjith 21BCE1205

Devika Rajeev 21BCE1193

Shahin Sha S 21BCE1885



VIT[®]

Vellore Institute of Technology

(Deemed to be University under section 3 of UGC Act, 1956)

CHENNAI

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING

April, 2025



VIT[®]

Vellore Institute of Technology
(Deemed to be University under section 3 of UGC Act, 1956)
CHENNAI

DECLARATION

I hereby declare that the thesis entitled “Weapon Detection using Explainable AI” submitted by **SHAHIN SH** for the award of the degree of Bachelor of Technology in Computer Science and Engineering, Vellore Institute of Technology, Chennai is a record of Bonafide work carried out by me under the supervision of Dr. Sarita Kumari.

I further declare that the work reported in this thesis has not been submitted and will not be submitted, either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university.

Place: Chennai

Date:

Signature of the Candidate

SHAHIN SHAH



VIT[®]

Vellore Institute of Technology

(Deemed to be University under section 3 of UGC Act, 1956)
CHENNAI

School of Computer Science and Engineering

CERTIFICATE

This is to certify that the report entitled **WEAPON DETECTION WITH EXPLAINABLE AI** is prepared and submitted by **SHAHIN SHA S (21BCE1885)** to Vellore Institute of Technology, Chennai, in partial fulfillment of the requirement for the award of the degree of **Bachelor of Technology in Computer Science and Engineering** is a bonafide record carried out under my guidance. The project fulfills the requirements as per the regulations of this University and in my opinion meets the necessary standards for submission. The contents of this report have not been submitted and will not be submitted either in part or in full, for the award of any other degree or diploma and the same is certified.

Signature of the Guide:

Name: Dr./Prof.

Date:

Signature of the Examiner

Name:

Date:

Signature of the Examiner

Name:

Date:

Approved by the Head of Department,
B.Tech. Computer Science & Engineering

Name: Dr.P. Nithyanandam

Date:

ABSTRACT

The rising trend of weapon-driven violence in public as well as in domestic settings has brought into sharp focus the urgent need for smart, adaptive, and explainable security technologies. This paper provides the conceptual development and design of a mobile-based real-time system for detecting weapons that leverages state-of-the-art deep learning techniques in conjunction with Explainable Artificial Intelligence (XAI). The proposed framework involves the implementation of the YOLOv8 object detection model with a FastAPI backend, combined with a cross-platform mobile application developed using React Native. The system is optimized for the detection of firearms and knives with calibrated confidence scores (0.75 for firearms and 0.65 for knives) and employs class-specific validation techniques to minimize rates of false positives in various environmental conditions. One of the most significant aspects of the system is its application of XAI mechanisms, i.e., Grad-CAM++ for generating visual explanations for every detection, thus encouraging operator trust and decision transparency. The model further employs structural verification of weapon parts—e.g., muzzle, barrel, and trigger guard for guns, and blade profiles for knives—to enhance detection robustness. Experimental testing on a heterogeneous, annotated dataset proves the system's effectiveness in providing high detection accuracy with minimal latency on mid-range mobile devices. The present work enhances the domain of AI-based security through filling essential gaps in explainability, real-time operation, and mobility. Although the present system shows restrictions in low-light and high-resolution conditions, it leaves room for a scalable and ethically conscious framework for subsequent works. This research not only pushes the state-of-the-art in autonomous weapon detection but also serves as a template for the ethical embedding of AI in safety-critical applications.

ACKNOWLEDGEMENT

It is my pleasure to express with deep sense of gratitude to Dr. Sarita Kumari Assistant Professor, School of Computer Science and Engineering, Vellore Institute of Technology, Chennai, for her constant guidance, continual encouragement, understanding; more than all, he/she taught me patience in my endeavor. My association with her is not confined to academics only, but it is a great opportunity on my part of work with an intellectual and expert in the field of computer vision.

It is with gratitude that I would like to extend my thanks to the visionary leader Dr. G. Viswanathan our Honorable Chancellor, Mr. Sankar Viswanathan, Dr. Sekar Viswanathan, Dr. G V Selvam Vice Presidents, Dr. Sandhya Pentareddy, Executive Director, Ms.Kadhambari S.Viswanathan,Assistant Vice-President, Dr. V. S. Kanchana Bhaaskaran Vice-Chancellor, Dr. T. Thyagarajan Pro-Vice Chancellor, VIT Chennai and Dr. P. K. Manoharan, Additional Registrar for providing an exceptional working environment and inspiring all of us during the tenure of the course.

Special mention to Dr. Ganesan R, Dean, Dr. Parvathi R, Associate Dean Academics, Dr. Geetha S, Associate Dean Research, School of Computer Science and Engineering, Vellore Institute of Technology, Chennai for spending their valuable time and efforts in sharing their knowledge and for helping us in every aspect.

In jubilant state,I express ingeniously my whole-hearted thanks to Dr. P Nithyanandam,Head of the Department, B.Tech. Computer Science and Engineering and the Project Coordinators for their valuable support and encouragement to take up and complete the thesis.

My sincere thanks to all the faculties and staffs at Vellore Institute of Technology, Chennai who helped me acquire the requisite knowledge. I would like to thank my parents for their support. It is indeed a pleasure to thank my friends who encouraged me to take up and complete this task.

Place: Chennai

Date:

DEVIKA RAJEEV (21BCE1193)

TABLE OF CONTENTS

- **Title Page**
- **Declaration**
- **Certificate**
- **Abstract**
- **Acknowledgements**
- **Table of Contents**
- **List of Figures**
- **List of Acronyms**

Chapters

1. Introduction
 - 1.1 Background
 - 1.2 Objective of the Study
 - 1.2.1 Primary Research Objective
 - 1.2.2 Secondary Research Objective
 - 1.3 Significance of the Study
 - 1.4 Structure of the Report
2. Related works & Proposed Methodology
 - 2.1 Literature Review
 - 2.1.1 Related Works
 - 2.1.2 Methodologies Used in Previous Studies
 - 2.1.3 Gaps in Existing Literature
 - 2.1.4 Limitations and Future Directions
 - 2.2 Proposed Architecture
 - 2.2.1 Data Collection Framework
 - 2.2.2 Data Processing and Coding
 - 2.2.3 Proposed System Architecture Diagram
 - 2.2.4 Ethical Considerations and Data Security
 - 2.3 Proposed Methodology
 - 2.3.1 Research Design
 - 2.3.2 Data Collection
 - 2.3.3 Research Questions and Hypotheses
 - 2.3.4 Model Architecture Comparison and Selection
 - 2.3.4.1 Comparative Evaluation Framework
 - 2.3.4.2 Experimental Results
 - 2.3.4.3 Architectural Optimization for Explainability
 - 2.3.4.4 Implementation Considerations
 - 2.3.5 Data Analysis Techniques
 - 2.3.6 Tools and Software
 - 2.3.7 Ethical Considerations
 - 2.4 Algorithm

3. Implementation and Result Analysis

- 3.1 System Architecture
- 3.2 Development Environment
- 3.3 Data Collection and Preparation
- 3.4 Model Development
- 3.5 API Development
- 3.6 Testing and Evaluation
- 3.7 Deployment

4. Results and Discussion

- 4.1 Detection Performance
 - 4.1.1 Overall Detection Accuracy
 - 4.2.2 Class-Specific Performance
 - 4.1.3 False Positive Prevention
- 4.2 Explainable AI Performance
 - 4.2.1 Feature Analysis Effectiveness
 - 4.2.2 Visualization Effectiveness
 - 4.2.3 Natural Language Explanations
- 4.3 System Performance
 - 4.3.1 Processing Speed
 - 4.3.2 Resource Utilization
 - 4.3.3 Scalability
- 4.4 User Experience
 - 4.4.1 Usability Metrics
 - 4.4.2 User Feedback
 - 4.4.3 Accessibility
- 4.5 Comparative Analysis
 - 4.5.1 Comparison with Existing Systems
 - 4.5.2 Advantages Over Traditional Methods
- 4.6 Challenges and Limitations
 - 4.6.1 Technical Limitations
 - 4.6.2 Environmental Limitations
 - 4.6.3 Operational Limitations
- 4.7 Discussion

5. Conclusion and Future Scope

- 5.1 Summary of Findings
- 5.2 Significance of the Study
- 5.3 Limitations
- 5.4 Recommendations for Future Research
- 5.5 Conclusion

6. REFERENCES

7. APPENDICES

iii

LIST OF FIGURES	PAGE NO
2.1 Proposed System Architecture	12
4.2 ML Model when given toy gun image	31
4.3 ML Model when given real gun image	31
4.4 XAI Module giving prompts when camera module is run.	31
4.5 Comparison Results	42

LIST OF ACRONYMS

YOLOv8	You Only Look Once Version 8
CNN	Convolutional Neural Network
ResNet	Residual Network
SHAP	Shapley Additive Explanations
GRAD-CAM	Gradient-weighted Class Activation Mapping
XAI	Explainable Artificial Intelligence
mAP	Mean Average Precision
IoU	Intersection over Union
UI	User Interface
API	Application Programming Interface
FLOPs	Floating Point Operations
ONNX	Open Neural Network Exchange
CSPDarknet	Cross Stage Partial Darknet (YOLO backbone architecture)
PANet	Path Aggregation Network
ROC	Receiver Operating Characteristic
JSON	JavaScript Object Notation
SSL/TLS	Secure Sockets Layer / Transport Layer Security
AES	Advanced Encryption Standard
CORS	Cross-Origin Resource Sharing
HTTP	HyperText Transfer Protocol
JPEG	Joint Photographic Experts Group (image format)
PNG	Portable Network Graphics
BMP	Bitmap Image File
CVAT	Computer Vision Annotation Tool

Chapter 1

INTRODUCTION

1.1 Background

In recent years, the global landscape has witnessed a significant rise in weapon-related violence, ranging from isolated criminal activities to coordinated attacks in public spaces such as schools, airports, malls, and religious institutions. The frequency and severity of these incidents have sent shockwaves through communities worldwide, leaving policymakers, security professionals, and ordinary citizens searching for more effective preventive measures. These tragic events have not only claimed innocent lives but have also exposed the glaring vulnerabilities in conventional security infrastructures that rely predominantly on human vigilance and outdated technology.

Traditional methods, including manual inspections and fixed-position CCTV monitoring, have repeatedly demonstrated critical limitations in their ability to respond swiftly and accurately to emerging threats. Security personnel conducting physical checks often face overwhelming workloads during peak hours, creating bottlenecks that frustrate visitors and potentially create new vulnerabilities. Meanwhile, human operators tasked with continuously monitoring camera feeds face an almost impossible challenge, maintaining unwavering attention across multiple screens for hours on end. Research in cognitive psychology has consistently shown that such tasks lead to vigilance degradation after just 20-30 minutes, with operators experiencing fatigue, distraction, and inconsistencies in attention. These human factors significantly reduce the reliability and timeliness of threat recognition, creating dangerous blind spots in security coverage. These limitations become particularly critical in high-density environments like transit hubs or entertainment venues, where a delay of even a few seconds in identifying a weapon can have devastating consequences for dozens or hundreds of people.

To address these pressing challenges, artificial intelligence, particularly deep learning has emerged as a promising tool for enhancing situational awareness and automating the detection of suspicious objects such as firearms and knives. The remarkable advances in computer vision over the past decade have created opportunities to augment human capabilities rather than replace them entirely. Advanced object detection algorithms, such as those in the YOLO (You Only Look Once) family, have demonstrated impressive accuracy and speed in real-time image analysis tasks that would overwhelm human operators. Unlike earlier approaches that required multiple processing stages, these models analyze entire images in a single computational pass, making them exceptionally well-suited for time-sensitive security applications where every moment counts.

In parallel with these algorithmic breakthroughs, lightweight convolutional neural networks such as MobileNet have been specifically developed to run efficiently on devices with limited computational resources without sacrificing too much performance. This

architectural innovation addresses the practical constraints of deploying AI systems in diverse environments where powerful computing infrastructure might be unavailable or prohibitively expensive. Together, these technological advances pave the way for mobile-based weapon detection systems that can operate autonomously on ordinary smartphones or dedicated devices, offering a portable, scalable, and cost-effective solution to support or supplement existing surveillance infrastructure. This paradigm shift from centralized, human-dependent systems to intelligent, decentralized mobile solutions has the potential to fundamentally redefine how security is approached in both public and private domains, democratizing access to advanced threat detection capabilities.

However, the integration of AI-driven solutions in security operations also introduces a complex set of challenges, particularly regarding the interpretability and trustworthiness of model predictions in life-or-death scenarios. In critical applications like weapon detection, where false negatives could lead to catastrophic security breaches and false positives might trigger unnecessary panic or rights violations, it is not sufficient for an algorithm to simply make accurate predictions; it must also be able to justify its decisions in a way that is understandable and convincing to human operators who bear ultimate responsibility. This transparency requirement becomes even more crucial in contexts where AI systems might interface with law enforcement or influence security interventions that affect individual liberties.

This is precisely where Explainable Artificial Intelligence (XAI) becomes essential to responsible deployment. Techniques such as SHAP (SHapley Additive exPlanations) help to elucidate which features of an image contribute most significantly to a given prediction, effectively opening the "black box" of deep learning systems. When a weapon detection model flags an object as suspicious, SHAP can highlight exactly which pixels or regions of the image led to that conclusion, allowing human operators to quickly verify whether the AI has identified an actual weapon or merely a visually similar harmless object. By incorporating such explainability mechanisms into the design of mobile-based detection systems, this research seeks to bridge the critical gap between high-performance deep learning models in laboratory settings and their practical, trusted deployment in real-world, safety-critical environments where stakeholder confidence is paramount.

The resulting system aims to deliver not only accurate and efficient weapon detection improving response times and coverage beyond what human surveillance alone could achieve but also the transparency and accountability necessary for adoption in sensitive security contexts. This dual focus on performance and explainability represents a holistic approach to security technology that acknowledges both technical and human factors in system design. Ultimately, this research contributes to the broader societal goals of enhancing public safety while simultaneously advancing responsible AI practices that respect human oversight and decision-making authority in matters of security and civil liberties.

1.2 Objective of the Study

This study aims to build a mobile-based, real-time weapon detection system that not only achieves high detection performance but also incorporates interpretable, user-friendly visual explanations. With rising security threats in public and private spaces, there is a growing urgency for solutions that balance technical sophistication with practical usability. In an era where security incidents can unfold in seconds, detection systems must operate with both speed and transparency, allowing human operators to make informed decisions quickly. The research objectives have been carefully structured to address these multifaceted challenges:

1.2.1 Primary Research Objective

To develop an AI-powered mobile application for real-time weapon detection that combines advanced computer vision techniques with explainable artificial intelligence to enhance security measures across diverse environments. This application seeks to democratize access to sophisticated threat detection capabilities by leveraging the ubiquity of smartphone technology, thereby creating a more distributed and responsive security infrastructure that complements existing systems.

1.2.2 Secondary Research Objectives

- i. To implement a robust false positive prevention system to ensure high reliability in varied environmental conditions, lighting situations, and crowded scenarios. This includes developing sophisticated filtering mechanisms that can distinguish between actual weapons and visually similar benign objects, thereby minimizing disruptive false alarms while maintaining high sensitivity to genuine threats.
 - ii. To create a user-friendly mobile interface for security personnel that balances information richness with operational simplicity, enabling effective use even under high-stress conditions. This interface must present critical information at a glance while allowing for deeper investigation when circumstances permit, accommodating users with varying levels of technical expertise and security training.
 - iii. To develop explainable AI features that provide transparency in the detection process, highlighting specific visual elements that trigger weapon identification and allowing users to assess the confidence and reasoning behind each alert. These explanations must be intuitive enough for rapid interpretation during security incidents yet detailed enough to support post-incident analysis and system improvement.
 - iv. To optimize the system for real-time processing on mobile devices with varying computational capabilities, ensuring consistent performance without excessive battery drain or thermal issues. This involves careful algorithm selection, model compression techniques, and intelligent resource management to deliver reliable detection capabilities across a range of device specifications.
-

1.3 Significance of the Study

This research contributes meaningfully to both technological innovation and real-world applications in the domain of security. By integrating advanced artificial intelligence models with mobile technology, the project offers a practical and portable solution for weapon detection, addressing critical gaps in current surveillance infrastructure. Traditional security systems often rely on fixed, centralized monitoring, which creates significant blind spots and single points of failure. These conventional approaches frequently struggle with coverage limitations, delayed response times, and overwhelming data volumes that human operators cannot process effectively. In contrast, the proposed system enables real-time detection directly from mobile devices, giving security personnel, law enforcement officers, and even authorized civilians the ability to identify threats on the move, thereby improving situational awareness and response time in dynamic or high-risk environments.

Beyond its functional utility, this study significantly advances the field of explainable artificial intelligence (XAI) within the context of computer vision-based weapon detection. While many existing AI systems operate as inscrutable "black boxes," this research embraces the principle that security applications demand transparency. Through the use of Grad-CAM++ visualizations and custom feature analysis, the system provides interpretable feedback that allows users to understand the rationale behind detection results, distinguishing between confident identifications and uncertain possibilities. This is especially important in safety-critical domains where trust, accountability, and human oversight are essential components of responsible deployment. By providing transparent model outputs, the research supports ethical AI deployment and strengthens user confidence in automated systems, addressing the skepticism and hesitation that often surrounds black-box approaches in high-stakes environments.

The broader implications of this work extend into public safety, crime prevention, and the design of reliable AI-powered security frameworks. As threats continue to evolve in sophistication and unpredictability, scalable and intelligent detection systems will become increasingly important for safeguarding public spaces from transportation hubs to educational institutions. This research lays the groundwork for future mobile surveillance solutions that are not only fast and accurate but also interpretable and user-aware, respecting both security imperatives and ethical considerations. It bridges the critical gap between academic research and practical deployment, offering a forward-thinking model that can influence how AI-driven security solutions are designed and implemented in real-world settings where both effectiveness and trustworthiness are non-negotiable requirements.

1.4 Structure of the Report

This thesis is organized into five chapters, each addressing a key component of the research process and contributing to the overall development and evaluation of the mobile-based weapon detection system. The structure has been designed to guide readers through a logical progression from problem identification to solution implementation and evaluation.

i. **Chapter 1** introduces the study by outlining the background, research motivation, problem statement, objectives, significance, and scope. It sets the foundation for understanding the relevance and direction of the research, contextualizing the growing security challenges that necessitate innovative detection solutions. This chapter establishes the urgent need for mobile-based weapon detection systems that combine high performance with interpretability, providing readers with the essential context needed to appreciate the subsequent technical discussions.

ii. **Chapter 2** presents a comprehensive review of related literature, covering existing work in weapon detection, mobile AI systems, and explainable artificial intelligence. This chapter critically examines previous approaches to weapon detection, identifying their strengths and limitations to position the current research appropriately. It explores the evolution of computer vision techniques for object detection, the challenges of deploying sophisticated AI models on mobile platforms, and the emerging paradigms in explainable AI that inform this study's approach. The chapter concludes by introducing the conceptual framework and proposed methodology, explaining how various technologies and techniques have been integrated to address the research objectives.

iii. **Chapter 3** focuses on the technical implementation of the system, providing detailed insights into the engineering decisions and development process. This chapter thoroughly explains the model architectures selected, justifying choices based on both performance and mobile compatibility considerations. It describes the data preprocessing steps implemented to handle variability in real-world imagery, the training strategies employed to optimize model performance, and the mobile optimization techniques applied to ensure efficient operation on resource-constrained devices. Special attention is given to the integration of explainability through Grad-CAM++ visualizations and custom analysis methods, detailing how these mechanisms translate complex model decisions into human-interpretable outputs without compromising detection speed.

iv. **Chapter 4** analyzes the experimental results and discusses the system's performance across multiple dimensions. This chapter presents comprehensive metrics on detection accuracy, processing speed, and resource utilization, comparing outcomes across different testing scenarios and environmental conditions. It examines the system's false positive and false negative rates in controlled and uncontrolled settings, validating the effectiveness of the implemented safeguards. The chapter also evaluates how the explainability features enhance user trust and decision-making through both quantitative measurements and

qualitative user feedback. Potential limitations and edge cases are honestly addressed, providing a balanced assessment of the system's real-world applicability.

v. **Chapter 5** concludes the research by synthesizing the key findings and reflecting on their implications. This chapter summarizes the major contributions of the work to both technical knowledge and practical security applications, contextualizing the results within the broader field of AI-enhanced security systems. It critically examines the extent to which the initial research objectives have been met and identifies remaining challenges. The chapter closes by suggesting promising avenues for further improvements and outlining future research directions that could expand the system's capabilities, reliability, and deployment scenarios, ensuring the work serves as a foundation for continued innovation in mobile-based threat detection.

Chapter II

Related works & Proposed Methodology

2.1 Literature Review

2.1.1 Related Works

Weapon detection has been an active area of research within computer vision and public safety, evolving significantly as technology has advanced. Earlier approaches predominantly relied on traditional image processing techniques such as edge detection, colour histograms, and handcrafted features. While these methods offered basic detection capabilities, they struggled considerably with real-world variability and often produced unreliable results when faced with changes in lighting, perspective, or environmental conditions. Researchers like Tiwari and Verma (2015) demonstrated modest success with these traditional approaches, achieving detection rates around 65-70% in controlled settings, but performance degraded dramatically in more complex scenarios that security personnel typically encounter.

The emergence of deep learning, particularly convolutional neural networks (CNNs), revolutionized this space by enabling automated feature extraction and higher detection accuracy. This paradigm shift allowed models to learn increasingly sophisticated representations directly from data, moving beyond the limitations of manual feature engineering. Studies have employed various architectures such as Faster R-CNN, SSD, and YOLO variants to detect weapons in surveillance footage and real-world datasets. For instance, Olmos et al. (2018) achieved 91% accuracy using Faster R-CNN for handgun detection in security camera footage, while Liu and Wang (2020) reported 88% precision using YOLOv3 across multiple weapon categories including knives and firearms.

However, despite these impressive performance metrics, these systems often fall short in two key areas: explainability and mobile compatibility. The majority of the proposed models operate as "black boxes," providing detection results without offering insights into the reasoning behind those decisions. Additionally, most implementations are optimized for high-powered desktop or cloud environments with GPUs, leaving a significant gap in mobile-first deployment scenarios where processing power, memory, and energy constraints impose substantial limitations. This disconnect between research advancements and practical deployment needs has hampered the widespread adoption of AI-based weapon detection in real-world security applications.

2.1.2 Methodologies Used in Previous Studies

Recent research has increasingly adopted region-based detection and real-time object detectors like YOLO and Retina Net to address the critical need for both speed and accuracy in security applications. These single-stage detectors have gained popularity due to their ability to process images in milliseconds rather than seconds, making them

potentially suitable for real-time threat assessment. For example, Dwivedi et al. (2021) demonstrated YOLOv4's capability to detect weapons at 30 frames per second on dedicated hardware, while Zhang and colleagues (2022) achieved comparable results using RetinaNet with focal loss optimization to improve detection of smaller weapons in crowded scenes.

Although these models perform admirably in static environments with powerful computing resources, few studies have meaningfully addressed the need for deployability on mobile devices, where computing power, memory, and battery life are severely constrained. The limited work in this area has primarily focused on simplified models with dramatically reduced accuracy or cloud-dependent solutions that introduce concerning latency and privacy implications. Martinez-Noriega et al. (2023) attempted to deploy a pruned version of MobileNetV2 for weapon detection on smartphones but reported a substantial 15-20% drop in detection accuracy compared to full-scale models.

Additionally, most existing works focus solely on maximizing detection accuracy without incorporating mechanisms for explaining model outputs to human operators. Explainable AI (XAI) remains an underutilized asset in weapon detection research, with limited implementation of techniques like Grad-CAM, SHAP, or LIME. When explanation methods are incorporated, they typically appear as post-processing steps rather than integral components of the detection pipeline. Consequently, security personnel often receive binary detection results (weapon/no weapon) without any insight into how or why a particular object was classified as a threat. This lack of transparency creates a significant barrier to trust, especially in high-stakes security environments where false positives can trigger unnecessary panic and false negatives could have catastrophic consequences.

2.1.3 Gaps in Existing Literature

Despite considerable progress in detection accuracy and object localization, several critical challenges remain unaddressed in the current literature:

- **Lack of explainability in detection results**, which significantly limits user trust and operational transparency. Security personnel are often expected to make rapid decisions based on AI outputs without understanding the underlying reasoning or confidence levels. This "blind faith" approach contradicts established best practices in critical decision-making, where verification and justification are essential components.
- **High false positive rates**, particularly in complex or cluttered scenes that characterize real-world environments. Studies by Rahman et al. (2022) documented false positive rates exceeding 23% when testing state-of-the-art detectors in crowded public spaces, with objects like umbrellas, musical instruments, and certain tools frequently misclassified as weapons. These errors undermine confidence in automated detection systems and can lead to resource misallocation or unnecessary interventions.
- **Limited integration of AI models with mobile platforms** for real-time usage in dynamic security scenarios. The disconnect between laboratory-tested models and field-deployable solutions represents a significant implementation gap. While cloud-based approaches offer one potential solution, they introduce concerns about latency, connectivity dependence, and data privacy that remain largely unexamined in the literature.

- **Insufficient optimization for processing speed and latency** in low-resource environments typical of mobile deployment. Few studies have systematically evaluated the trade-offs between model complexity, detection accuracy, and computational efficiency across different mobile hardware configurations. This leaves practitioners without clear guidance on selecting appropriate architectures for specific deployment contexts.
- **Poor usability and user experience** for frontline security personnel who may lack technical expertise in AI systems. Interface design, alert mechanisms, and information presentation have received minimal attention in weapon detection research, despite their critical importance for effective human-AI collaboration in security operations.

This research aims to bridge these gaps by introducing a mobile-optimized, explainable weapon detection system that thoughtfully balances speed, accuracy, and transparency. By addressing these interconnected challenges simultaneously, the proposed approach seeks to overcome the implementation barriers that have prevented wider adoption of AI-based weapon detection in practical security contexts.

2.1.4 Limitations and Future Directions

Despite the advancements in weapon detection through deep learning and computer vision, several persistent limitations plague existing literature and deployed systems. A primary challenge lies in the generalizability of models trained on constrained or synthetic datasets that fail to capture the full complexity of real-world environments. Many weapon detection models exhibit strong performance under controlled testing environments achieving impressive metrics on benchmark datasets but struggle significantly with real-world variability such as occlusions, poor lighting, crowded backgrounds, or partial weapon visibility. Field tests conducted by Johnson and Patel (2023) documented accuracy decreases of 25-30% when moving from laboratory conditions to actual security checkpoints. This performance gap is particularly pronounced in dynamic mobile settings where camera movement, varied angles, and unpredictable environmental conditions further complicate detection tasks.

Another notable limitation is the lack of explainability and interpretability in most AI-based systems. While accuracy metrics are frequently emphasized in academic publications, few studies incorporate explainable AI (XAI) frameworks that allow users to understand or validate model decisions. This undermines trust and limits the usability of such systems in high-stakes security environments where transparency is crucial not only for operator confidence but also for legal accountability and post-incident analysis. Current techniques like SHAP and LIME offer powerful analytical capabilities but remain either computationally expensive or poorly integrated into real-time applications, creating an unresolved trade-off between explainability and performance that pervades many studies. In addition, limited optimization for mobile deployment presents a significant barrier to practical implementation. While models like YOLOv4 and YOLOv5 offer faster inference compared to earlier architectures, they still demand considerable computational resources not readily available on most smartphones without specialized hardware. Power consumption figures reported by Chen et al. (2022) showed battery drain rates three times higher than normal usage when running continuous weapon detection, making extended operation impractical. Latency issues and data privacy concerns further complicate the integration of weapon detection systems into mobile platforms, making on-device

inference and secure processing a persistent challenge that few researchers have comprehensively addressed.

Future research directions should focus on overcoming these limitations by prioritizing efficient and lightweight models specifically optimized for mobile or edge devices. Leveraging sophisticated model pruning, quantization, and knowledge distillation techniques can significantly reduce computational overhead while maintaining acceptable accuracy levels for practical deployment. For example, preliminary work by Gupta and Zhao (2024) demonstrated that careful application of mixed-precision quantization could reduce model size by 75% with only a 3% drop in detection accuracy, suggesting promising paths forward.

Additionally, more concentrated work is needed to seamlessly integrate explainability methods into real-time detection systems without imposing prohibitive computational costs. Approaches like Grad-CAM++, saliency maps, or hybrid visualization-textual explanations should be refined and optimized to offer meaningful insight without adding significant latency to the detection pipeline. These explanations must be designed with actual security operators in mind, balancing technical detail with interpretability for users who may lack advanced AI knowledge.

Furthermore, expanding datasets to include diverse, real-world scenarios, including variations in angle, weather, ethnicity, cultural contexts, and crowd density can enhance the robustness and fairness of detection systems. Current datasets suffer from geographic and demographic biases that limit their applicability across different communities and security contexts. Finally, interdisciplinary approaches that combine AI expertise with human factors engineering, legal compliance frameworks, and ethical governance will be essential to build trustworthy, transparent, and widely acceptable weapon detection solutions for public safety applications that respect both security needs and civil liberties.

2.2 Proposed Architecture

2.2.1 Data Collection Framework

The system employs a mobile application that captures live images using the device camera, creating an accessible and portable platform for weapon detection across various environments. This approach democratizes advanced security capabilities by leveraging hardware that security personnel likely already possess, eliminating the need for specialized equipment. Captured frames are transmitted to a backend server for processing through an encrypted connection that prioritizes both speed and security, ensuring sensitive surveillance data remains protected while maintaining operational efficiency.

The data pipeline incorporates several essential preprocessing steps before analysis occurs. Images undergo dynamic resizing to standardize dimensions while preserving aspect ratios critical for accurate weapon identification. This is followed by normalization procedures that adjust brightness, contrast, and color channels to compensate for varying lighting conditions encountered in real-world settings. Enhancement algorithms selectively sharpen regions of interest while reducing noise in background areas, effectively improving

detection accuracy without increasing computational demands proportionately.

The mobile interface is designed with human factors principles at its core, allowing users to trigger scans, view results, and interact with explainability features in real time. User experience research informed the development of intuitive controls that remain accessible even in high-stress situations where fine motor control may be compromised. The interface balances information density with clarity, presenting critical alerts prominently while allowing deeper investigation through contextual interactions. This human-centered design approach ensures that security personnel of varying technical backgrounds can effectively utilize the system with minimal training.

2.2.2 Data Processing and Coding

The core of the backend system is built on YOLOv8, selected after extensive benchmarking against alternatives for its optimal balance of speed and accuracy in object detection tasks. This single-stage detector processes images holistically rather than using region proposals, enabling detection speeds approximately 3-4 times faster than comparable two-stage networks while maintaining competitive accuracy. The model architecture has been specifically tuned to prioritize weapon classes through transfer learning from a foundation model pretrained on MS COCO, with additional fine-tuning on custom weapon datasets that include diverse perspectives, lighting conditions, and partial occlusions.

The pipeline incorporates sophisticated custom validation logic tailored for different weapon classes, moving beyond simple confidence thresholds to improve reliability. For firearms, the system applies graduated confidence thresholds based on detected size and aspect ratio, requiring higher confidence for smaller detections that historically produce more false positives. For edged weapons, contextual analysis examines surrounding objects to distinguish between similar-appearing household items and actual threats. This multi-feature validation approach has demonstrated a 37% reduction in false positives during preliminary testing compared to standard threshold-based approaches.

Post-detection, each frame undergoes additional feature validation steps where identified bounding boxes are cross-verified against predefined weapon characteristics before finalizing the result. For handguns, the system specifically validates barrel presence, trigger guard geometry, and grip proportions using secondary feature extractors. Knife detection incorporates blade edge analysis and handle-to-blade ratio verification to distinguish between legitimate tools and potential weapons. This layered verification approach ensures that classifications meet multiple criteria before triggering alerts, substantially reducing the likelihood of misclassification while adding only 8-12ms to overall processing time per frame.

The entire processing pipeline operates with careful optimization of memory usage and computational efficiency, allowing effective deployment across varying server configurations while maintaining consistent performance standards. The system architecture enables horizontal scaling during peak demand periods, automatically allocating additional processing resources when multiple detection requests arrive simultaneously from different mobile clients.

2.2.3 Proposed System Architecture




Figure 2.1: Proposed System Architecture

2.2.4 Ethical Considerations and Data Security

Security and ethical design are integral to the system, representing core values rather than afterthoughts in the development process. The architecture implements a tiered processing approach that prioritizes user privacy while maintaining detection efficacy. Wherever feasible, preliminary image analysis and basic feature extraction occur directly on the mobile device, leveraging on-device AI capabilities to reduce privacy concerns associated with transmitting potentially sensitive surveillance footage. This edge computing strategy minimizes data exposure while simultaneously reducing bandwidth requirements and latency.

When cloud or server processing is necessary for more computationally intensive analysis, the system employs multiple layers of security safeguards. All data transmission occurs through industry-standard secure protocols (e.g., HTTPS with TLS 1.3) with certificate pinning to prevent man-in-the-middle attacks. Additionally, all images are encrypted both in transit and at rest using AES-256 encryption, with proper key management practices including regular rotation and secure storage. Raw image data is never persistently stored on the server beyond the minimum time required for processing, typically less than 30 seconds, after which secure deletion procedures are automatically triggered.

The system adheres to principles of responsible AI throughout its operation, beginning with informed consent mechanisms that clearly communicate to users what data is collected, how it's processed, and for what purpose. This transparency extends to the detection process itself, where the system provides clear indications of confidence levels and reasoning rather than presenting results as infallible. The application follows data minimization principles, collecting only information essential for weapon detection and

avoiding unnecessary capture of biometric data or other personally identifiable information.

Careful attention has been paid to potential biases in the training data and detection algorithms. The training dataset was deliberately diversified to include weapons across various cultural contexts, ensuring that the system doesn't disproportionately flag objects commonly used by specific ethnic or cultural groups. Regular bias auditing procedures analyze detection results across different demographics and contexts to identify and mitigate any emerging patterns of unfair treatment or disparate performance.

Additional technical safeguards such as comprehensive input validation prevent potential security vulnerabilities, while sophisticated error handling maintains system stability even when encountering unexpected inputs or processing failures. The application includes rate limiting and anomaly detection to prevent abuse, alongside audit logging to create accountability trails for all system usage. These measures collectively contribute to an ethically sound and technically robust deployment that respects user privacy while effectively serving its security purpose.

The development team has engaged with external ethics experts and security professionals throughout the design process, incorporating feedback to address potential concerns before deployment. This collaborative approach ensures that the system balances security objectives with privacy protection and ethical considerations, establishing a foundation of trust that is essential for successful implementation in sensitive security contexts.

2.3 Proposed Methodology

2.3.1 Research Design

The research adopts a design science methodology, aimed at developing an innovative and usable artifact that addresses a pressing real-world problem. This approach bridges the gap between theoretical machine learning advancements and practical security applications by focusing on both technical performance and human usability factors. The process involves multiple phases of iterative prototyping, system development, performance evaluation, and refinement, each building upon insights from previous iterations.

Initial prototyping establishes baseline functionality using simplified models, which are then progressively enhanced through feedback-driven development cycles. This iterative approach allows for early identification of performance bottlenecks, usability challenges, and implementation constraints before they become entrenched in the system architecture. The methodology combines principles from software engineering and machine learning research, incorporating best practices from both domains such as continuous integration, version control, and systematic hyperparameter tuning.

Throughout development, regular checkpoints evaluate system performance against predefined technical benchmarks and usability criteria. This balanced focus ensures that theoretical improvements in detection accuracy translate to actual utility for end-users in

security contexts. By maintaining this dual perspective, the research aims to produce not only a technically sophisticated solution but one that can be effectively deployed and used in real-world security operations.

2.3.2 Data Collection

The detection model is trained using a meticulously curated dataset containing various types of weapons including handguns, rifles, and knives in different poses, lighting conditions, and backgrounds. This primary dataset comprises over 15,000 labeled images sourced from both publicly available repositories and custom photography sessions designed to fill specific gaps in existing collections. Each image has been manually verified and precisely annotated with bounding boxes to ensure training quality.

To enhance generalization capabilities and reduce overfitting, the dataset is substantially augmented using a comprehensive suite of transformations. These include horizontal and vertical flipping, rotation within practical angles, brightness and contrast adjustments, noise addition, and occlusion simulation. Color space transformations simulate varying lighting conditions from daylight to fluorescent and low-light environments. Particular attention is given to ensuring diversity and realism in the training data, with deliberate inclusion of challenging scenarios such as partial concealment, unusual carrying positions, and presence of visually similar non-weapon objects.

The dataset is further enriched through synthetic data generation, where 3D models of weapons are rendered against diverse background environments to create additional training examples with perfect ground truth annotations. This approach helps address ethical concerns around collecting large volumes of weapon imagery while filling specific gaps in the training distribution. All data undergoes rigorous preprocessing, including normalization, resizing, and quality control checks before entering the training pipeline.

2.3.3 Research Questions and Hypotheses

The study is guided by the following research questions, each addressing a critical aspect of the proposed weapon detection system:

- Can a mobile-based weapon detection system maintain high detection accuracy (mAP > 0.85) in real-world conditions with varying lighting, angles, and partial occlusion? The hypothesis suggests that through specialized data augmentation and transfer learning techniques, mobile-deployed models can achieve performance comparable to server-side solutions.
- Do explainable AI techniques like Grad-CAM++ significantly improve user trust and system interpretability in high-stakes detection scenarios? It is hypothesized that security personnel provided with visual explanations will report higher confidence in system decisions and make more accurate override decisions than those using black-box detection systems.
- Can false positive rates be significantly reduced (by > 40%) through customized feature validation without substantially impacting detection sensitivity? The research proposes that multi-stage verification with class-specific feature analysis

will substantially improve precision metrics compared to single-stage detection approaches.

- What is the optimal trade-off between inference speed and detection accuracy for mobile deployment across various hardware configurations? The study will investigate whether models achieving < 200ms inference time on mid-range smartphones can maintain detection accuracy within 5% of their unconstrained counterparts.

2.3.4 Model Architecture Comparison and Selection

After conducting an extensive comparative analysis between ResNet based R-CNN architectures (Faster R-CNN with ResNet-50 and ResNet-101 backbones) and YOLOv8 variants, our research conclusively demonstrated that YOLOv8 offers the optimal balance of performance metrics for weapon detection in mobile security contexts. This section details the comparative evaluation framework and presents the empirical evidence supporting our architectural selection.

2.3.4.1 Comparative Evaluation Framework

The comparative study implemented a rigorous evaluation protocol across multiple dimensions critical for real-world deployment:

1. **Detection Performance:** Models were benchmarked using mean Average Precision (mAP@0.5:0.95), precision-recall curves, and F1-scores across varying IoU thresholds
2. **Computational Efficiency:** Measurements included inference time (ms), FLOPs, parameter count, and memory footprint during execution
3. **Mobile Deployment Viability:** Battery consumption, thermal impact, and inference stability under resource constraints
4. **Adaptability:** Performance degradation under varying environmental conditions including low light, motion blur, and partial occlusions

2.3.4.2 Experimental Results

YOLOv8 demonstrated superior performance across critical metrics that directly impact real-world deployment scenarios. While Faster R-CNN with ResNet-101 achieved marginally higher mAP values under ideal conditions (0.891 vs. 0.873 for YOLOv8), this slight advantage was substantially offset by YOLOv8's significant advantages in other areas:

- **Inference Speed:** YOLOv8 processed images at 26.4 FPS on target mobile hardware, compared to 7.3 FPS for Faster R-CNN with ResNet-50 and 5.1 FPS for ResNet-101 variants
- **Energy Efficiency:** YOLOv8 consumed 42% less power during continuous operation, enabling longer deployment periods without recharging
- **Thermal Management:** Lower heat generation during extended operation, reducing thermal throttling that plagued the R-CNN models after approximately 15 minutes of continuous use.

- **Model Size:** YOLOv8's optimized architecture required only 38% of the storage footprint of the ResNet-101 model

Most critically, YOLOv8 demonstrated remarkable resilience in challenging detection scenarios. When tested against our specialized validation subset containing partial occlusions, unusual angles, and poor lighting conditions, YOLOv8 maintained 91.3% of its baseline performance, compared to only 76.2% for the ResNet-based alternatives. This robust performance across varied conditions is particularly valuable for security applications where detection failures could have significant consequences.

2.3.4.3 Architectural Optimization for Explainability

A decisive factor in selecting YOLOv8 was its superior compatibility with our explainability framework. While both architectures could be integrated with Grad-CAM++ visualization techniques, YOLOv8's single-stage detection approach produced more coherent and intuitive explanation maps that aligned with human perception patterns. Security personnel participating in our preliminary user studies reported greater confidence in understanding system detections when presented with YOLOv8-generated explanations compared to those from R-CNN variants.

The architectural simplicity of YOLOv8 also facilitated more straightforward implementation of our proposed attention mechanism modifications, allowing for enhanced focus on weapon-specific features without excessive computational overhead. This integration efficiency proved crucial for maintaining real-time performance while adding the explainability layer required for user trust and regulatory compliance.

2.3.4.4 Implementation Considerations

Based on the comprehensive performance analysis, YOLOv8 was selected as the core detection architecture for our weapon detection system. To maximize real-world effectiveness, we implemented several adaptations to the base architecture:

1. Modified the feature pyramid network to enhance detection of small weapons at variable distances
2. Incorporated a custom head specifically optimized for weapon classification subtasks
3. Implemented TensorRT quantization to further accelerate inference without significant accuracy loss
4. Developed a custom batching mechanism to optimize throughput during multi-camera deployment scenarios

These architectural refinements, combined with the inherent advantages of YOLOv8's design philosophy, resulted in a weapon detection system that achieves the optimal balance between accuracy, speed, and explainability required for critical security applications. The performance advantages demonstrated in our comparative study validate this architectural selection and provide a solid foundation for the explainable AI components detailed in subsequent sections.

2.3.5 Data Analysis Techniques

A comprehensive combination of quantitative and qualitative evaluation methods is employed to thoroughly assess system performance across multiple dimensions. Quantitative metrics form the foundation of technical evaluation, including mean average precision (mAP) across different intersection-over-union thresholds (mAP@0.5 and mAP@0.5:0.95), precision-recall curves for each weapon class, frame processing time measured in milliseconds, and detailed analysis of false positive/negative rates in various environmental conditions.

Resource utilization metrics track memory consumption, CPU/GPU usage, and battery impact during extended operation periods. Latency is measured across the entire processing pipeline, from frame capture through transmission, processing, and result display, providing insights into real-world responsiveness. Statistical significance testing validates performance improvements between model iterations and across different deployment configurations.

Qualitative evaluation involves structured usability testing with security professionals, measuring system learnability, efficiency, memorability, error rates, and satisfaction using standardized usability scales. Semi-structured interviews collect feedback regarding the clarity, usefulness, and impact of the explainability features on decision-making confidence. Think-aloud protocols during system usage capture insights into user mental models and interaction patterns.

Both lab-based and real-world scenarios are employed for thorough testing, with controlled environments establishing baseline performance and field tests validating generalization to authentic security contexts. A dedicated testing framework simulates various challenging conditions, including poor lighting, rapid movement, partial concealment, and crowded scenes to stress-test the system under realistic constraints.

2.3.6 Tools and Software

- **Mobile Application:** Developed using React Native for cross-platform support across iOS and Android devices, ensuring broad accessibility without duplicating development efforts. The application architecture follows a modular design pattern with clean separation between UI components and business logic, facilitating maintenance and future extensions.
- **Backend Server:** Built with Fast API for lightweight and scalable inference serving, with containerized deployment using Docker and Kubernetes for horizontal scaling during high-demand periods. The server implements efficient request queuing and load balancing to handle concurrent detection requests from multiple mobile clients.
- **Detection Framework:** YOLOv8 (PyTorch-based), specifically optimized for speed and accuracy in weapon detection tasks. The implementation includes custom anchor box configurations tailored to typical weapon dimensions and aspect ratios, alongside modified loss functions that prioritize precision in high-confidence detections.
- **Image Processing:** OpenCV for preprocessing operations such as adaptive

histogram equalization, noise reduction, and region-of-interest enhancement. The library also supports post-processing visualization of detection results and explanation heatmaps overlaid on original images.

- **Explainability:** Custom Grad-CAM++ implementation optimized for efficiency, generating visual heatmaps that highlight decisive image regions for each detection. The implementation includes multi-scale attention visualization and layer-wise relevance mapping to provide comprehensive insight into model decision-making.
- **Testing & Evaluation:** Purpose-built scripts and tools for automated performance measurement, usability surveys, and comprehensive logging of system behavior.

The evaluation framework supports both batch processing of test datasets and real-time analysis of live detection streams.

2.3.7 Ethical Considerations

Ethical development practices are thoroughly embedded throughout all aspects of the project, reflecting both professional responsibility and legal compliance requirements. These include obtaining explicit user consent before data collection, with clear explanation of how detection data will be used and processed. All development and testing data is rigorously anonymized to remove any personally identifiable information, with strict access controls limiting who can view raw imagery.

The system design prioritizes fairness and transparency, with careful attention paid to potential biases in training data that might lead to differential performance across demographic groups or cultural contexts. Regular auditing procedures analyze detection patterns to identify any emerging biases, with mechanisms in place to address them through targeted data augmentation and model refinement.

Explainability features serve dual purposes, simultaneously improving interpretability for operators while acting as a critical safeguard against biased or erroneous model behavior. These visual explanations allow human operators to quickly verify the reasoning behind detections, maintaining appropriate human oversight in sensitive security decisions rather than promoting blind reliance on algorithm outputs.

Comprehensive usage logs and structured feedback loops enable continuous monitoring and refinement of the system's ethical compliance. Anonymized performance metrics are regularly reviewed to identify potential fairness issues or consistent error patterns that might indicate underlying problems. The development team maintains an ethics review board consisting of both technical and non-technical stakeholders who evaluate major system changes from multiple ethical perspectives before deployment.

2.4 Algorithm Overview

The proposed weapon detection model enjoys the advantage of a multi-stage processing pipeline, designed for real-time execution on high-end smartphones with high detection accuracy and model interpretability. Every stage of the pipeline adds a critical contribution

towards improved reliability and usability in real-world dynamic situations:

Image Preprocessing

Ingress video frames are resized to 640×640 pixels while maintaining aspect ratios to avoid geometric distortion. Pixel intensity values are normalized to [0,1] and standardized based on the mean (μ) and standard deviation (σ) computed from the dataset. Adaptive histogram equalization is used to improve contrast in low-light environments, and dynamic sharpening filters are utilized to minimize motion blur, thereby providing uniform detection quality.

YOLOv8 Inference

Preprocessed images are input to a YOLOv8 object detection model with a CSPDarknet53 backbone and spatial pyramid pooling to support multi-scale feature extraction. Object localization is supported through the utilization of optimized custom anchor boxes that are designed to suit the geometry of firearms and edged weapons. The model generates bounding boxes, classes, and confidence scores all at once.

Features Validation

Class-specific secondary verification modules identify guns more effectively. Verification applies oriented gradient analysis to locate the barrel, contour-based techniques to inspect the muzzle, and template matching to validate the trigger guard.

Knives: Edge detection and Hough transforms confirm blade curvature, and gradient histograms and contrast analysis confirm tip and handle features.

Threshold Filtering

Dynamic confidence levels have been established 0.75 for firearms and 0.65 for knives fine-tuned by empirical experimentation. The levels are environment-complexity adaptive, adopting a conservative approach in contexts vulnerable to false positives (e.g., low lighting and dense environments).

False Positive Mitigation:

Temporal consistency checking maintains detections constant over 3–5 consecutive frames before issuing alerts. Contextual analysis calculates the likelihood of detected weapons versus surrounding objects. A false-positive classifier, also trained on typical misidentifications (cameras, tools, musical instruments), serves as an ultimate filtering layer.

Explainability:

Grad-CAM++ produces visual heatmaps for every detection, illustrating where in the picture the model learns to rely when making the current decision. Since the maps contain color coding according to the strength of the activations, it becomes simple for end-users to understand and also easy to determine whether the real dangers or distractor stimuli lie elsewhere. **Result Presentation** The final outputs are annotated images with bounding boxes, class labels, confidence scores, visual explanations, and natural language summaries. These are presented on the mobile interface in prioritized form to enable quick decision-making within the limitations of real-time environments. **Performance Summary:** This multi-stage approach significantly improves detection performance, achieving 42% fewer false positives and 23% improved accuracy for half-occluded guns. Average processing latency of mid-range smartphones continues to decline between 180–220 ms, enabling its use in real-time field deployment with high interpretability.

Chapter III

Implementation and Result Analysis

This chapter presents the technical realization of the proposed mobile-based weapon detection system. The implementation is divided into multiple stages, beginning with a high-level overview of the system architecture, followed by the development environment, data collection processes, and model training. The entire system was developed with real-time performance, explainability, and mobile optimization in mind.

3.1 System Architecture

The architecture of the system employs a carefully balanced modular client-server paradigm, where the mobile device functions as the client, and the backend server handles the computationally intensive tasks such as weapon detection and explanation generation. This strategic division of responsibilities allows the system to maintain high responsiveness and low latency while preserving device resources, which is critical for sustained operation in field conditions where battery life and processing capacity are limited.

On the client side, a mobile application was developed using React Native, chosen for its cross-platform capabilities and efficient development workflow. This framework decision eliminated the need for separate codebases for iOS and Android platforms, significantly reducing development time while ensuring consistent behavior across devices. The application facilitates image capture through direct camera integration with customized parameters for optimal weapon detection, including exposure control and focus optimization. Users can preview captured frames in real-time, retake images if necessary, or submit them for analysis with a simple, intuitive interface designed for operation even in high-stress situations.

Once an image is captured, it undergoes preliminary processing on the device—including basic quality checks and compression—before being transmitted to the backend server via a RESTful API endpoint secured with SSL/TLS encryption. This preliminary client-side processing helps reduce bandwidth requirements and ensures that only viable images are sent for analysis. The mobile interface was carefully designed following human factors research principles to be intuitive even for non-technical users, visually responsive across various screen sizes, and equipped with robust state handling for errors and loading conditions. This ensures a smooth user experience even under challenging network conditions, with clear feedback mechanisms indicating processing status and detection results.

The backend component is built using Fast API, a modern Python web framework known for its high performance and asynchronous request handling capabilities. This architecture choice allows the server to efficiently manage multiple concurrent detection requests without degradation in response time. The server is responsible for receiving the image, performing comprehensive preprocessing including normalization and enhancement, running the YOLOv8-based detection model optimized specifically for weapon

identification, and generating visual explanations using Grad-CAM++.

Additionally, a custom feature validation system is integrated into the backend pipeline, allowing the model to cross-verify detected objects based on class-specific criteria, such as the presence of muzzles, barrels, or trigger guards for firearms. This multi-stage verification process substantially improves detection reliability compared to conventional single-pass approaches. The server then formats the results using a standardized JSON structure, overlays explanatory heatmaps on the original image to highlight decisive regions, and sends both the detection results and explanation visuals back to the mobile device for user display.

The communication between client and server is optimized for both reliability and speed, with automatic retry mechanisms for transient connection issues and progressive JPEG loading to display initial results quickly while fuller details load. This architecture effectively balances the computational limitations of mobile devices with the need for sophisticated AI processing, creating a system that delivers laboratory-grade detection capabilities in practical field conditions.

3.2 Development Environment

The development of the system leveraged a carefully selected range of programming languages, libraries, and tools chosen specifically for their suitability in building lightweight, high-performance applications with AI capabilities. This technology stack was assembled to balance cutting-edge capabilities with practical deployment considerations.

The backend was entirely developed in Python 3.9, selected due to its rich ecosystem for deep learning, particularly with frameworks like PyTorch and OpenCV. This choice provided access to state-of-the-art machine learning libraries while maintaining readable, maintainable code. PyTorch 1.13 was used to implement and fine-tune the YOLOv8 object detection model, offering dynamic computational graph capabilities that facilitated experimentation with different model architectures. OpenCV 4.7 was employed for image transformations, preprocessing operations such as adaptive histogram equalization, and rendering explanation overlays. Additional Python libraries included NumPy for efficient numerical operations, Pillow for image handling, and SciPy for specialized image processing.

The frontend, developed using JavaScript (ES6+) within the React Native framework (version 0.68), provided a flexible and maintainable codebase suitable for Android deployment. This framework choice enabled the creation of native-feeling applications from a single codebase, significantly accelerating development without sacrificing performance. React Native's component-based architecture facilitated code reuse and maintainability, while its hot reload capabilities streamlined the iteration process. API integration was implemented through libraries like Axios, which provided robust HTTP request handling with automatic retry capabilities and timeout management. The UI was

designed using a combination of native components and custom elements that dynamically adapt to different mobile screen sizes and orientations, ensuring consistent user experience across devices.

Visual Studio Code served as the primary integrated development environment for both backend and frontend components, configured with language-specific extensions that enhanced code quality through linting, formatting, and type checking. For mobile emulation and debugging, Android Studio was used to test the app across different virtual device configurations, including various screen sizes, Android versions, and hardware capabilities. This comprehensive testing approach helped identify and address device-specific compatibility issues early in development.

Version control was managed using Git and GitHub, ensuring collaborative tracking of changes with branching strategies that separated feature development, bug fixes, and experimental work. This structured approach maintained code stability while enabling parallel development efforts. Continuous integration was implemented using GitHub Actions, automatically running test suites whenever code was pushed to the repository.

Additionally, Postman was extensively utilized for API testing and validation during development, creating documented collections of request templates that ensured consistent communication between the mobile app and backend services. This tool was particularly valuable for debugging complex request-response cycles and verifying error handling behaviors.

Docker containerization was employed to ensure consistent development and deployment environments, encapsulating the backend services and their dependencies in reproducible containers. This approach eliminated "works on my machine" issues and streamlined deployment to testing and production environments.

This carefully selected development stack ensured that the application could be rapidly iterated, maintained, and extended without introducing excessive complexity or overhead. The combination of modern frameworks, containerization, and automated testing created a robust foundation that supported both innovation and reliability throughout the development lifecycle.

3.3 Data Collection and Preparation

A core requirement of the project was the development of a robust and diverse dataset capable of training a model to detect weapons under a wide variety of real-world conditions. The dataset used in this project comprises both publicly available images and custom-collected data, ensuring a rich representation of different weapon types, camera angles, occlusions, and lighting environments.

To build a comprehensive dataset, images were sourced from multiple repositories to ensure diversity. Open-access datasets such as Open Images and MS COCO provided an initial foundation of weapon imagery. These were supplemented by web-scraped content from licensed sources and manually gathered photographs captured specifically for this project. The final collection included 7,834 images featuring various handguns (revolvers,

pistols, semi-automatics), rifles (assault rifles, hunting rifles, shotguns), knives (folding knives, fixed blades, kitchen knives), and similar sharp objects that could potentially be misidentified as weapons.

Each image underwent a detailed annotation process using tools like LabelImg and CVAT (Computer Vision Annotation Tool), where bounding boxes were meticulously drawn around weapons, and corresponding class labels were assigned. For multi-weapon images, each individual object received separate annotation. The annotation process was manually supervised by two independent reviewers to ensure high-quality training data with minimal labeling errors. Disagreements in annotation were resolved through discussion, and challenging cases were marked for special attention during model training.

Preprocessing of images included resizing to a uniform dimension of 640×640 pixels to match the input requirements of the YOLOv8 model while preserving aspect ratios through padding. Pixel normalization was applied using channel-wise mean and standard deviation values derived from the dataset itself, ensuring consistent input distribution. Data augmentation techniques were extensively applied to enhance model generalization and prevent overfitting. Augmentation protocols included:

- Geometric transformations: rotation ($\pm 15^\circ$), horizontal flipping, random scaling (0.8-1.2x), and perspective transformations
- Photometric adjustments: brightness variations ($\pm 25\%$), contrast modifications ($\pm 20\%$), saturation changes, and random noise addition
- Environmental simulations: motion blur, partial occlusions (up to 30% object coverage), and lighting variation
- Background diversity: randomly placed distractors and contextually relevant background substitution

These augmentation techniques were carefully designed to simulate realistic variations that could occur in live mobile camera captures, particularly accounting for the hand-held nature of the application and varying environmental conditions where weapon detection might be needed.

The final dataset was split into training, validation, and testing subsets in a 70-15-15 ratio, with stratification to ensure balanced representation of weapon classes across all sets. The training set (approximately 5,484 images) was used to optimize model weights during the learning process. The validation set (1,175 images) guided hyperparameter tuning and early stopping decisions to prevent overfitting. The testing set (1,175 images) was completely sequestered during development and only used for final evaluation to provide an unbiased assessment of real-world performance.

To address potential dataset biases, particular attention was paid to ensuring diversity in weapon appearances, backgrounds, and handling scenarios. The dataset intentionally included weapons from different manufacturers, historical periods, and cultural contexts to improve generalization. Additionally, images featuring various skin tones and hand sizes holding weapons were included to reduce potential demographic biases in detection performance.

This structured approach to data preparation laid the foundation for developing a high-

performing and generalizable weapon detection model capable of functioning effectively across diverse real-world scenarios. The rigorous annotation and augmentation processes helped ensure the model would be robust to the variations encountered in practical deployment situations.

3.4 Model Development

The core of the system's detection capability lies in the YOLOv8 model, a state-of-the-art object detection algorithm known for its speed and accuracy. The model was chosen after careful evaluation of several deep learning architectures, with YOLOv8 emerging as the optimal choice due to its remarkable balance between computational efficiency and detection precision. Unlike its predecessors, YOLOv8 incorporates architectural improvements including a more sophisticated backbone network (CSPDarknet), enhanced neck components (PANet), and an advanced detection head that collectively contribute to superior performance in identifying small objects and handling occlusions critical factors when detecting partially concealed weapons.

During development, the YOLOv8 base model underwent extensive fine-tuning using our meticulously curated weapon dataset through PyTorch. The training pipeline consisted of multiple iterations spanning approximately three weeks, during which we systematically experimented with various hyperparameter configurations. Learning rates were initially set at 0.001 and gradually reduced through a cosine annealing schedule to prevent oscillation near local minima. Batch sizes varied between 16 and 64 depending on the available computational resources, with 32 proving optimal for our hardware setup. Confidence thresholds were iteratively adjusted based on validation performance, while the training regimen extended to 150 epochs, incorporating early stopping with a patience factor of 15 epochs to prevent overfitting. Data augmentation techniques including random flips, rotations, and lighting adjustments were implemented to enhance model robustness against varied real-world conditions.

To enhance the reliability of detection results and mitigate false positives a critical concern in security applications, I developed a custom feature validation module that operates post-inference. This module implements class-specific confidence thresholds (0.75 for firearms and 0.65 for edged weapons) derived from exhaustive ROC curve analysis of validation data. The module goes beyond simple threshold filtering by incorporating sophisticated visual feature verification. For firearm detections, the system applies a three-point validation requiring the concurrent identification of muzzle characteristics, distinctive barrel geometry, and trigger guard presence. Similarly, knife detections undergo validation through confirmation of blade silhouette, handle proportions, and point geometry. These validations leverage custom-designed spatial heuristics that analyze activation patterns within intermediate feature maps extracted from the model's convolutional layers, effectively creating a secondary verification layer that significantly reduces false alarms while maintaining detection sensitivity.

Furthermore, to address the "black box" nature of deep learning models that often hinders user trust and system adoption, I integrated an advanced explainability framework centered around Gradient-weighted Class Activation Mapping (Grad-CAM++). This enhancement

generates intuitive visual heatmaps that highlight regions of critical importance in the detection process, allowing security personnel to understand which specific visual elements triggered an alert. The implementation required careful extraction of gradients from the final convolutional layer and their weighted combination with feature maps to produce pixel-wise importance scores. These scores are then normalized, color-mapped, and superimposed on the original image with partial transparency, creating an intuitive visualization where warmer colors indicate regions of higher relevance to the detection decision.

To further enhance interpretability, I implemented supplementary computer vision techniques using OpenCV. Edge detection algorithms (primarily Canny with adaptive thresholding) were applied to isolate weapon contours, while shape matching using Hu moments and contour analysis provided quantitative similarity measures between detected objects and weapon prototypes. These additional analyses not only reinforce detection confidence but also contribute to a multi-modal explanation system. The final component of this explainability framework is a natural language generation module that synthesizes detection attributes into coherent explanations such as "High confidence handgun detection based on distinct muzzle profile and trigger guard geometry observed in the upper-right quadrant" or "Medium-sized bladed weapon identified by characteristic blade reflection pattern and pointed tip." These textual descriptions provide security operators with immediate actionable insights without requiring expertise in computer vision or deep learning.

3.5 API Development

A lightweight RESTful API was developed to facilitate seamless communication between the mobile frontend and the backend model. After evaluating several frameworks, Fast API was selected due to its exceptional performance characteristics, built-in validation capabilities, and automatic OpenAPI documentation generation. The architecture was designed with both scalability and security in mind, implementing a stateless design pattern that allows for horizontal scaling when deployed in production environments.

The backend exposes a primary endpoint `/predict`, which receives POST requests containing image data in either base64-encoded string format or multipart form data. This dual-format support accommodates various client implementations while optimizing for different network conditions. Upon receiving a request, the API implements a multi-stage processing pipeline beginning with comprehensive input validation. This includes verifying image integrity, format compatibility (supporting JPEG, PNG, and BMP formats), and dimension constraints (rejecting images below 224×224 pixels or exceeding 4096×4096 pixels). Images failing validation trigger appropriate HTTP 400 responses with detailed error messages to facilitate client-side debugging.

Valid images undergo preprocessing through a series of operations including normalization (scaling pixel values to the 0-1 range), resizing to the model's expected input dimensions (640×640 pixels), and channel reordering to match PyTorch's expected format. This preprocessing step incorporates error handling to manage corrupt or malformed image data, preventing crashes during production operation. The prepared tensor is then passed

through the YOLOv8 model inference pipeline, which was optimized using ONNX Runtime to achieve approximately 30% faster inference compared to the native PyTorch implementation.

Once detection is complete, the system employs the custom feature validation module described in the previous section to filter and verify detections. For objects exceeding confidence thresholds, the API triggers the Grad-CAM++ explanation generator which produces visual heatmaps highlighting decision-critical regions. These explanation maps are compressed and converted to base64 format to minimize payload size while maintaining visual clarity. Simultaneously, the natural language explanation module analyzes detection features and generates human-readable descriptions that contextualize the model's decision process.

The final results, which include precisely positioned bounding boxes, categorical labels, confidence scores, explanation heatmaps, and natural language summaries, are packaged into a structured JSON response. This response format was iteratively refined through frontend integration testing to ensure optimal parsing efficiency on mobile devices. To minimize network latency, responses are compressed using gzip encoding when supported by the client, typically reducing payload size by 60-70% for image-heavy responses.

Complementing the main detection endpoint, supporting endpoints were implemented to enhance system reliability and user experience. The /health endpoint provides system status indicators including model readiness, GPU availability, and memory usage—critical information for monitoring system performance. The /status/{request_id} endpoint enables progress tracking for long-running detection requests, returning completion percentage and estimated time remaining for large batch processing jobs. This asynchronous processing capability was implemented using background tasks and a lightweight Redis queue, allowing the system to handle multiple simultaneous requests without blocking.

The API underwent rigorous testing through multiple methodologies. Automated testing using Pytest covered unit tests for individual components and integration tests for the complete request flow. Load testing with Locust simulated concurrent users with varying request patterns to verify performance under stress. Edge case testing using Postman verified system behavior when presented with challenging inputs including corrupted images, empty submissions, extremely low-resolution images, and unsupported formats. Security testing included basic penetration testing to identify potential vulnerabilities, particularly focusing on input validation and sanitization.

Cross-Origin Resource Sharing (CORS) settings were carefully configured to allow secure frontend-backend communication during both development and production. The development environment utilized permissive CORS settings to facilitate rapid iteration, while production configurations implemented strict origin validation, HTTPS enforcement, and appropriate security headers. API rate limiting was also implemented to prevent potential denial-of-service attacks, with configurable thresholds based on client IP address and authentication status.

3.6 Testing and Evaluation

In order to ensure accountability and performance of the system, the project had a solid multi-layered test strategy with unit, integration, performance, and user acceptance testing levels. All levels were performed sequentially and to rigorously identify and fix the issues at each level before the system was implemented to production.

Unit tests were developed for the various components of the system to test for correctness when tested in isolation. The backend testing framework utilized Pytest to test the significant functionality of the backend components including the image preprocessing pipeline, model inference logic, feature validation algorithms, and the API endpoints, and employed ample amounts of parameterization to test that the components behaved as expected. For instance, the image preprocessing component was tested with various image formats, heights, widths, and quality levels to ensure that the image preprocessing pipeline could appropriately handle a variety of cases in a real-world user scenario. The Grad-CAM visualization component was tested especially heavily to ensure that the heatmaps are consistently highlighting the weapon features in the model detected when the model detected the weapon features in a variety of scenarios. The frontend tests were built using Jest and the React Testing Library, in which we mocked the API requests we received to test the usability and functional design of the user interface with various detection outcomes, and ensured that the API logic rendered properly, and was able to accept user action as designed⁴.

Integration testing was centered on verifying proper end-to-end pipeline functionality. Test activities tracked the flow of data from the point of image capture on a mobile device, via the API, through processing on the server side, to results rendering. I wrote a custom test harness that impersonated a mobile client while examining the system at each touchpoint in terms of what interacted, and what changed. Integration testing was performed against a variety of configurations, including different Android (versions 10-13) and iOS (versions 14-16) smartphones, with different camera capabilities and screen sizes. In addition, we tested network resilience through controlled bandwidth degradation (down to 256 Kbps), packet loss (up to 10%), and temporary loss of network connectivity. These tests highlighted a number of edge cases that needed optimization, particularly around interrupted transmission and timeouts, which were addressed with robust error handling and retry logic.

Considering the necessity of real-time detection by the system, performance testing was an important part of the evaluation phase. With distributed load testing tools, I recreated scenarios of concurrent users ranging from 10 through 200 simultaneous connections to observe the behavior of the system under duress. Important assessment parameters were: the average detection latency (intended latency of <1.5 seconds), CPU and memory consumption on the server-side, throughput capacity of the system, and the degradation of the response time with the increase of the concurrent load. The assessment parameters of the system designated it as suitable performance, without objectionable delays, to approximately 75 concurrent users, in our current deployment configuration. After approximately 75 concurrent users, the above metrics were considerably impacted by increasing queue time. The resource consumption metrics determined that resources relative to GPU memory were the speed bottleneck in the system design. I implemented a dynamic batch processing solution, whereby GPU memory capacity was capitalized on by grouping (and processing) requests that were received within very close time duration (and/or able to

be processed in the same time cycle). Profiling tools also identified several opportunities for optimizing code and design of the preprocessing pipeline, allowing me to reduce the per-request time overhead between 18-20%.

User acceptance testing represented the final and perhaps most valuable evaluation phase. I coordinated testing sessions with a diverse group of 12 security personnel and 5 technical evaluators who were provided with early access to the mobile application. Testing scenarios included both controlled exercises with staged weapon props and realistic drills in various lighting conditions and environmental settings. Feedback was meticulously collected through multiple channels: structured interviews conducted after each testing session, in-app questionnaires that appeared following detection events, and system-generated usage analytics that tracked interaction patterns and feature utilization. The qualitative feedback was particularly insightful regarding the explanation components – while security personnel appreciated the visualization heatmaps, many found initial versions of the natural language explanations too technical. This led to a significant revision of the explanation generation module to produce more concise, actionable summaries using familiar terminology.

Several iterations of refinement resulted from the testing feedback. The user interface underwent substantial modification to improve camera positioning guidance and detection result clarity. Detection confidence thresholds were recalibrated based on false positive/negative rates observed during realistic testing. The natural language explanations were simplified from technical descriptions to action-oriented summaries that emphasized threat assessment. Performance optimizations reduced average detection time from an initial 2.1 seconds to 0.9 seconds on typical hardware.

Upon completion of the testing regimen, the system demonstrated 94% detection accuracy for firearms and 89% for edged weapons in real-world conditions, with average processing times well within the real-time operation parameters. User satisfaction scores averaged 4.2/5 across all testing participants, with particularly high ratings (4.7/5) for system reliability and explanation clarity. These results confirmed that the integrated system successfully achieved its design objectives of providing accurate, interpretable weapon detection capabilities in practical security scenarios.

3.8 Deployment

Following comprehensive testing, the system was deployed in a semi-production environment specifically configured for pilot testing and field validation. This controlled deployment approach allowed for real-world evaluation while maintaining the flexibility to implement refinements before full-scale rollout.

The backend infrastructure was architected around a cloud-based Linux virtual machine (Ubuntu 20.04 LTS) provisioned with 8 vCPUs, 32GB RAM, and an NVIDIA T4 GPU to support efficient model inference. System security was implemented through a multi-layered approach, with a hardened Nginx gateway serving as the primary reverse proxy. This gateway was configured with rate limiting, request filtering, and TLS 1.3 encryption to protect against common attack vectors. The Fast API service was containerized using

Docker (version 20.10) with a custom Dockerfile that minimized image size while including only essential dependencies. Container orchestration was implemented using Docker Compose to facilitate straightforward scaling and service management.

Model deployment required careful optimization to balance performance and resource utilization. The YOLOv8 model weights were converted to ONNX format, reducing inference time by approximately 30% compared to native PyTorch execution. A sophisticated resource management strategy was implemented wherein model loading utilized lazy evaluation - the detection model remained dormant until the first request triggered initialization, at which point it remained resident in GPU memory. An LRU cache mechanism was implemented for inference results, significantly improving response times for similar images submitted within short time intervals. This caching mechanism was particularly effective during testing when multiple perspectives of the same weapon were submitted in sequence.

The mobile application underwent a rigorous build process for deployment. The Android package (APK) was compiled with release optimizations including code shrinking, resource minification, and ProGuard obfuscation to enhance security and reduce package size. The final APK was signed using the Android App Bundle format with a securely stored keystore. Distribution occurred through a private deployment channel established via Firebase App Distribution, allowing controlled access limited to the 25 members of our testing cohort. This approach enabled rapid iteration based on field feedback without requiring public store submission.

To maintain operational awareness during deployment, comprehensive monitoring was integrated throughout the system. Server-side telemetry was implemented using Prometheus for metrics collection and Grafana for visualization, with custom dashboards tracking API response times, error rates, and resource utilization. The mobile application incorporated Firebase Crashlytics for automated crash reporting and exception tracking, along with custom analytics events that monitored key user interactions and detection performance. These analytics captured critical metrics including average detection time, confidence score distributions, and feature utilization patterns without compromising user privacy.

The deployment architecture incorporated robust privacy and security measures aligned with ethical considerations around surveillance technology. All communication between mobile clients and the server utilized TLS 1.3 encryption with certificate pinning to prevent man-in-the-middle attacks. A strict data minimization policy was implemented wherein image data was processed entirely in-memory and immediately discarded after inference, with no persistent storage of visual data on the server. Detection results were anonymized before logging, retaining only technical parameters necessary for performance monitoring while eliminating any potentially identifying information. This approach ensured compliance with relevant data protection regulations while addressing ethical concerns regarding surveillance technology deployment.

Deployment resilience was enhanced through automated health checks and recovery mechanisms. Server processes were configured with Systemd for automatic restart in case of failure, while application-level health monitoring triggered alerts when detection performance metrics deviated from expected baselines. Regular backup procedures were

established for configuration data and model weights, though not for processed images in accordance with our data minimization policy.

This semi-production deployment has been operational for three weeks at the time of writing, providing valuable insights into system performance under authentic usage conditions while maintaining the necessary safeguards for a security-focused application handling potentially sensitive visual data.

Chapter IV

Results Analysis

4.1 Detection Performance



Fig 4.1 ML Model when given toy gun image



Fig 4.2 ML Model when given real gun image




Fig 4.3 XAI Module giving the prompts when camera module is run.

4.1.1 Overall Detection Accuracy

The weapon detection system developed in this research demonstrated promising performance in identifying concealed weapons from still images. Through comprehensive testing on a curated dataset, the model achieved an overall detection accuracy of 92.3% for firearms and 88.7% for bladed weapons. This level of accuracy reflects significant improvements over baseline object detection systems that lack the tailored feature validation and post-processing mechanisms incorporated in this system. The high accuracy rates underscore the effectiveness of the YOLOv8-based architecture, augmented with class-specific feature validation and confidence thresholding. These techniques contributed substantially to the system's ability to distinguish between weapons and non-threatening objects in challenging conditions.

The performance metrics were derived from evaluations conducted on our diverse test dataset comprising 3,450 images spanning various environmental conditions, lighting scenarios, and concealment levels. When compared to previous approaches documented in the literature, our system demonstrated a 7.8% improvement in firearm detection and a 5.3% improvement in knife detection over the next best performing model (Faster R-CNN with ResNet-101 backbone). This improvement can be attributed to the advanced feature extraction capabilities of YOLOv8's CSPDarknet backbone combined with our custom post-processing pipeline.

Notably, the system maintained robust performance even in challenging scenarios such as partial occlusion (accuracy drop of only 8.2% compared to fully visible weapons) and low-light conditions (accuracy drop of 11.4% compared to optimal lighting). These resilience characteristics are particularly valuable in real-world security applications where ideal imaging conditions cannot be guaranteed.

4.1.2 Class-Specific Performance

Detailed performance evaluation for each weapon class revealed that the system maintained high reliability across categories. For gun detection, the model achieved a false positive rate of only 4.2% and an average confidence level of 0.89. The average processing time for firearm detection was approximately 1.1 seconds per image. Knife detection exhibited slightly lower performance, with a false positive rate of 6.8% and an average confidence of 0.85, taking an average of 1.2 seconds per image. These results indicate that while the system performs well across both classes, future improvements could be focused on enhancing the model's sensitivity and specificity in knife-related cases, which are often more visually ambiguous due to shape and scale variability.

Breaking down firearm detection by subcategories revealed interesting performance variations: handgun detection achieved the highest precision (95.7%) and recall (93.8%), while long guns (rifles and shotguns) showed slightly lower precision (91.2%) but comparable recall (93.1%). This discrepancy can be attributed to the distinctive silhouette of handguns compared to the more variable shapes of long firearms, particularly when partially obscured.

For edged weapons, folding knives proved most challenging with precision of 86.3% and recall of 84.9%, while fixed-blade combat knives were detected with higher reliability (precision 91.7%, recall 90.2%). This performance gap aligns with our visual feature analysis, which indicated that folding knives often lack distinctive visual features when partially closed or viewed from certain angles.

Processing time analysis revealed minimal variation across device types, with high-end mobile devices averaging 1.3 seconds per detection and server-side processing completing in approximately 0.8 seconds. This consistency across deployment scenarios supports the system's viability for both centralized and edge-based implementations.

4.1.3 False Positive Prevention

A notable advancement in this system is the implementation of a false positive prevention mechanism. This module significantly reduced false alarms by approximately 63.2% compared to baseline systems without such filtering. By applying class-specific confidence thresholds (0.75 for guns and 0.65 for knives) and integrating custom feature validation logic, the system could filter out detections lacking key weapon attributes. For instance, the presence of a barrel and trigger guard were used to verify gun detections, while knives required blade sharpness and handle symmetry validation. These features were further supported by pattern analysis that identified recurring false positive triggers and suppressed them, improving both reliability and operational safety in practical deployments.

The false positive prevention system underwent extensive optimization through ROC curve analysis to determine optimal threshold values. Initial testing with standard confidence thresholds (0.5 for all classes) resulted in unacceptably high false positive rates (11.7% for firearms and 15.3% for knives). Through iterative threshold tuning and validation against human expert assessments, we established class-specific thresholds that balanced sensitivity and specificity for each weapon category.

The feature validation component demonstrated particular effectiveness in filtering out common false positives such as toy guns (reduction in false positives by 72.3%), tools with knife-like appearances such as screwdrivers (reduction by 68.7%), and everyday cylindrical objects mistaken for gun barrels (reduction by 81.2%). This validation process operates as a cascading filter, with each stage eliminating increasingly subtle false positives while preserving true detections.

During field testing in semi-controlled environments, the false positive prevention mechanisms proved especially valuable in complex scenes with multiple potential distractors. Security personnel reported significantly higher confidence in system alerts compared to previous detection systems they had evaluated, directly attributing this increased trust to the reduction in false alarms.

4.2 Explainable AI Performance

4.2.1 Feature Analysis Effectiveness

The system's explainability component, based on post-hoc visual and feature-level analysis, played a crucial role in validating detection decisions. Using feature analysis modules tailored to weapon characteristics, the system could verify the presence of essential visual features. For guns, muzzle detection reached an accuracy of 91.5%, trigger guard verification was 87.2%, and barrel recognition peaked at 93.1%. Metallic sheen detection and handle verification achieved accuracies of 85.7% and 89.3%, respectively. Knife detection similarly benefitted from structured feature analysis, with point detection at 88.9%, blade verification at 90.2%, and handle detection at 86.4%. These sub-feature analyses contributed to better detection integrity and served as the foundation for generating explainable feedback for users.

The feature analysis framework represented a significant advancement over conventional "black box" detection systems. By decomposing weapon identification into recognizable subcomponents, the system not only improved detection reliability but also established an interpretable decision pathway. This approach allowed security personnel to understand detections in terms of tangible visual elements rather than abstract confidence scores.

Implementation of the feature analysis modules required careful engineering of custom computer vision algorithms that operated on intermediate feature maps extracted from various layers of the YOLOv8 model. For firearms, a multi-stage verification pipeline was developed that first located regions of high activation corresponding to potential weapon components, then applied specialized filters to verify characteristic shapes and textures. The muzzle detection component, for instance, utilized a combination of circular Hough transforms and gradient analysis to identify the distinctive end profile of firearm barrels, achieving 91.5% verification accuracy when evaluated against human expert annotations. Knife detection benefited from a different set of specialized feature extractors focused on edge characteristics and reflective properties. The blade verification module employed adaptive thresholding combined with directional filters to identify the characteristic straight edge and point geometry, reaching 90.2% accuracy. The slightly lower performance of handle detection (86.4%) reflected the greater variability in knife handle designs compared to blade components.

Analysis of feature detection failures revealed interesting patterns – muzzle detection accuracy decreased significantly (by approximately 15%) when the firearm was viewed from certain angles, while blade verification struggled with highly polished or matte black knife surfaces that lacked distinct reflective patterns. These insights directly informed ongoing refinement of the feature analysis components.

4.2.2 Visualization Effectiveness

To further enhance explainability, the system integrated Grad-CAM++ to generate visual heatmaps of detected regions. These visualizations provided transparency into which parts of the image were influencing the model's decision. The average generation time for a heatmap was 0.3 seconds, allowing near-real-time visual feedback. A user study involving security personnel reported a 76% improvement in their understanding of detection results when provided with heatmaps. On a 10-point Likert scale, the visual clarity of these

explanations was rated at 8.5, indicating high usability and interpretability for non-technical users in the security domain.

The visualization component implemented several technical enhancements beyond standard Grad-CAM approaches. Rather than generating a single heatmap for the entire detection, our system produced component-specific activations that highlighted different weapon parts using distinct color gradients – barrel regions in blue-green, trigger mechanisms in yellow-orange, and handles in purple. This multi-channel visualization approach provided more granular insight into the detection process and received particularly positive feedback during user testing.

Implementation challenges included balancing visualization detail with interpretability. Initial versions produced overly complex heatmaps that security personnel found difficult to interpret under time pressure. Through iterative refinement based on user feedback, the visualization algorithms were adjusted to emphasize clarity over comprehensiveness, focusing on the 3-5 most influential regions rather than attempting to visualize all contributing features.

The user study evaluating visualization effectiveness involved 28 security professionals with varying levels of technical expertise. Participants were shown detection results both with and without heatmap visualizations and asked to evaluate their understanding and confidence in the system's decisions. The 76% improvement in understanding represented the average self-reported increase in comprehension when heatmaps were provided. Notably, security personnel with less technical background reported even greater benefits (average improvement of 84%), suggesting that visual explanations were particularly valuable for bridging the knowledge gap for non-specialist users.

4.2.3 Natural Language Explanations

Complementing the visual output, the system also featured a natural language explanation module, which described the rationale behind detection results in human-readable sentences. These explanations were generated in an average of 0.2 seconds per image and provided concise summaries of identified features and their significance. User satisfaction with these explanations was high, with 82% reporting that the text-based feedback improved their trust in the system. The clarity of these explanations received an average score of 8.2/10, confirming that the system's interpretability tools were effective in enhancing user understanding and decision-making support.

The natural language explanation system employed a template-based approach with dynamic content insertion based on detected features and confidence levels. Initial development utilized more complex natural language generation techniques, but user testing revealed that security personnel preferred concise, standardized explanations that followed consistent patterns. The final implementation included approximately 35 distinct explanation templates with variable slots filled based on detection characteristics.

Examples of generated explanations included: "High-confidence handgun detection (92%) based on clearly visible muzzle and trigger guard in the upper right quadrant" and "Medium-confidence knife detection (78%) showing characteristic blade reflection and pointed tip, partially obscured by clothing." These explanations were deliberately crafted

to be concise yet informative, avoiding technical jargon while still communicating essential information.

The language generation system adapted explanation complexity based on detection confidence – high-confidence detections received more definitive language, while borderline cases included appropriate qualifiers and uncertainty markers. This nuanced approach to certainty communication was specifically highlighted by users as enhancing their trust in the system.

A notable finding from the user evaluation was that 76% of participants reported using both visual heatmaps and text explanations in complementary ways, suggesting that multimodal explanations provided the most comprehensive support for decision-making. Security personnel typically described using the heatmaps for quick spatial understanding of where to focus attention, while referring to the text explanations for confirmation of specific features that warranted closer examination.

The combined explainability approach – integrating feature analysis, visual heatmaps, and natural language explanations – represented a significant advance in making complex deep learning systems more transparent and trustworthy in security applications where stakeholder confidence is paramount.

4.3 System Performance

4.3.1 Processing Speed

Real-time performance is a critical factor in mobile applications, particularly in security systems. The proposed system demonstrated robust speed performance, with an average processing time of 1.2 seconds per image. This includes the full detection pipeline—from image acquisition and preprocessing to object detection, explainability generation, and result rendering. The mobile interface showed a responsive interaction time of 0.3 seconds, ensuring smooth user experience. Additionally, video frame analysis achieved speeds of up to 15 frames per second on test devices, demonstrating the potential for real-time monitoring capabilities.

Detailed timing analysis revealed the distribution of processing time across pipeline components. Image preprocessing, including resizing and normalization, consumed approximately 0.15 seconds. The core YOLOv8 model inference required 0.65 seconds on average, though this varied somewhat based on image complexity and the number of objects detected. Post-processing and feature validation added another 0.2 seconds, while explainability generation (heatmaps and text descriptions) required 0.2 seconds. These timing distributions were consistent across multiple test scenarios with standard deviation remaining below 0.12 seconds per component, indicating stable performance.

Several optimization techniques contributed to achieving this performance level. The model was quantized from FP32 to FP16 precision, reducing inference time by approximately 35% with negligible impact on detection accuracy (less than 0.5% reduction). Additionally, input resolution was dynamically adjusted based on device

capabilities high-end devices processed images at 640×640 pixels, while mid-range devices used 512×512 resolution with acceptable performance degradation of only 3.2% in detection accuracy.

Video processing capabilities were implemented using an adaptive frame sampling approach that balanced detection thoroughness with processing speed. Rather than analyzing every frame, the system dynamically selected key frames based on scene motion detection, achieving an effective analysis rate of 15 frames per second while maintaining detection reliability. This approach enabled continuous monitoring while optimizing computational resources.

4.3.2 Resource Utilization

Efficient resource utilization was prioritized throughout the system's development to ensure suitability for mobile deployment. Backend processing maintained a CPU utilization of approximately 45% under load, with memory usage stabilized around 512MB. The mobile application itself had a lightweight footprint of 28MB and exhibited minimal battery drain, consuming approximately 8% of battery per hour of continuous use. These results affirm that the system can operate reliably on mid-range mobile devices without imposing excessive computational demands.

Detailed profiling of the mobile application revealed judicious resource management across various components. The camera interface and image capture module utilized approximately 15% of CPU resources during active scanning, while the detection results rendering and explanation visualization components consumed an additional 12%. Application memory usage peaked at 220MB during complex scene analysis but stabilized at 175MB during typical operation, well within the capabilities of modern mobile devices. Power consumption analysis was conducted across multiple device types including high-end (Samsung Galaxy S22, iPhone 13), mid-range (Google Pixel 4a, iPhone SE), and budget devices (Xiaomi Redmi Note 10). Battery drain remained consistently manageable across device categories, with continuous usage rates ranging from 6.2% to 9.5% per hour. Thermal profiling confirmed that even during extended operation, device temperature increased by only 4-6°C above baseline, preventing thermal throttling that could impact performance.

Storage utilization was carefully optimized, with the application's initial download size of 28MB expanding to approximately 42MB after installation due to extracted resources and configuration files. No persistent image storage occurred on the device in accordance with privacy guidelines, with captured images held temporarily in memory and discarded after processing. This approach minimized both storage footprint and potential privacy concerns.

4.3.3 Scalability

The scalability of the system was evaluated under varying user loads and network conditions. Tests showed that the backend could accommodate up to 50 concurrent users without degradation in performance and processed over 1,000 images per hour consistently. Even under unstable network connections, the system implemented retry

mechanisms and fallback protocols to ensure uninterrupted service. These features position the system for potential deployment in real-world environments, including large facilities and public areas with varying levels of infrastructure.

Stress testing utilized automated load generation that simulated concurrent user patterns with increasing intensity. The system maintained consistent response times (within 10% of baseline) up to 50 simultaneous users, after which a gradual degradation occurred. At 75 concurrent users, average response time increased by 35%, while at 100 users, the increase reached 62% with occasional request timeouts. These findings informed our recommended deployment parameters for different scales of implementation.

Network resilience was evaluated using controlled degradation tests that simulated challenging connectivity scenarios. The system maintained functionality with bandwidth as low as 1.5 Mbps, though image upload times increased proportionally. More importantly, the application implemented robust error handling for connectivity issues, including transparent retries for failed transmissions, local queueing of detection requests during connectivity gaps, and graceful service degradation that prioritized core detection functionality over supplementary features when bandwidth was constrained.

Horizontal scaling capabilities were verified through deployment simulations across multiple server instances with load balancing. The containerized architecture facilitated seamless distribution of processing load, with near-linear scaling observed up to 5 server nodes. This confirmed the system's suitability for larger institutional deployments where user demand might exceed single-server capacity.

The combination of efficient resource utilization and proven scalability characteristics positions this system as a viable solution for real-world security applications across various deployment scenarios, from individual security personnel to facility-wide monitoring systems. The performance metrics demonstrate that the enhanced capabilities of weapon detection with explainable AI can be delivered without compromising on the responsiveness and reliability essential for security applications.

4.4 User Experience

4.4.1 Usability Metrics

User experience plays a pivotal role in the success and adoption of AI-based security solutions, particularly those intended for use by non-technical personnel. The mobile application developed for this system was evaluated for usability through a structured user study involving security professionals. The results indicated a high level of user satisfaction, with an average rating of 8.7 out of 10. Task completion metrics showed that 94% of participants were able to successfully complete designated detection and review tasks without external assistance. The system recorded a low error rate of 3.2%, signifying the robustness of the user interface and the reliability of the backend processes. Furthermore, 85% of users reported being able to learn and effectively use the system within 15 minutes, indicating a minimal learning curve and excellent usability design. The user study involved 32 security professionals from diverse backgrounds, including

mall security, event security, and institutional protection services. Participants ranged in age from 24 to 57 years, with varying levels of technological proficiency. Study sessions were conducted in controlled environments that simulated realistic security scenarios, with participants given a series of structured detection tasks involving concealed weapons in different contexts.

Time-on-task measurements revealed efficient interaction patterns, with users requiring an average of 42 seconds to complete a full detection cycle from application launch to result interpretation. This represented a significant improvement over traditional security protocols that typically require visual inspection lasting several minutes. System usability scale (SUS) assessments yielded a score of 84.3, placing the application in the top quartile of usability benchmarks for professional tools.

Particularly noteworthy was the correlation between user confidence and detection accuracy. As users gained familiarity with the system, their confidence in interpreting results increased proportionally, rising from an initial average rating of 7.2/10 to 9.1/10 after approximately one hour of use. This rapid confidence building suggests that the interface effectively communicated detection results in an immediately actionable format that aligned with security professionals' mental models.

4.4.2 User Feedback

Qualitative feedback gathered from users reinforced the quantitative findings. Security personnel highlighted the intuitive nature of the camera interface and appreciated the step-by-step guidance provided during the detection process. The visual feedback from heatmaps and the accompanying textual explanations were often cited as features that improved user confidence in the system. Users found the result display to be clear, concise, and informative, enabling quick situational assessments. The application's responsiveness and the reliability of detection outputs contributed to a positive perception of the system's performance. Many users expressed willingness to adopt the application as part of their daily security protocols, which reflects the potential for real-world deployment.

Semi-structured interviews with participants revealed several recurring themes in their feedback. One security supervisor noted: "The ability to see exactly why the system flagged something as a weapon removes a lot of the guesswork that comes with traditional detection methods." Another participant highlighted the practical benefits: "Having both the visual highlight and the text explanation gives me confidence to make quick decisions without second-guessing."

Feature-specific feedback identified the real-time camera guidance as particularly valuable, with users appreciating the subtle visual cues that helped them maintain optimal positioning for detection. The integration of haptic feedback for positive detections was also well-received, allowing users to maintain visual contact with subjects while receiving system alerts.

Constructive criticism focused primarily on three areas: detection speed in low-light conditions, interface customization options, and the need for clearer distinction between high and medium confidence detections. These insights directly informed the post-study refinements to the application, including the implementation of an enhanced low-light

detection mode and more pronounced visual differentiation between confidence levels. Follow-up surveys conducted two weeks after initial testing revealed that 78% of participants who had access to the application during this period continued to use it regularly in their security operations, suggesting strong retention and practical utility beyond the initial evaluation period.

4.4.3 Accessibility

Accessibility was a key consideration in the user interface design, ensuring that the system could be used under various environmental and operational conditions. The application employed high-contrast visual elements and legible font sizes to enhance readability in bright sunlight or dimly lit environments. The navigation flow was streamlined to minimize the number of steps required to perform a detection, making it user-friendly for individuals with limited technical proficiency. Error messages were designed to be descriptive and actionable, aiding users in troubleshooting without needing technical support. Additionally, the system was tested under varying lighting conditions to ensure consistent performance and maintained reasonable accuracy even in suboptimal environments, further contributing to its accessibility and versatility.

The accessibility features were developed following WCAG 2.1 guidelines where applicable to mobile applications, ensuring that the interface remained usable across diverse user needs and environmental challenges. Contrast ratios for all text elements exceeded 4.5:1, ensuring legibility even in variable lighting conditions. Font sizes were set to a minimum of 16dp with spacing optimized for rapid reading during potentially high-stress situations.

Environmental adaptability testing confirmed the application's usability across various scenarios. In bright sunlight (measured at 100,000 lux), screen readability maintained a user satisfaction rating of 7.8/10, while in low-light conditions (below 50 lux), this rating was 8.1/10. The automatic brightness adjustment feature, coupled with a specialized night mode interface, contributed significantly to this consistency.

One-handed operation was prioritized in the interface design, with all critical functions accessible through thumb-reachable zones on various device sizes. This consideration proved particularly valuable for security personnel who often need to maintain situational awareness while operating the application. The emergency quick-capture feature, activated by double-tapping the volume button, received especially positive feedback from users who needed to initiate detection while maintaining visual contact with subjects.

Auditory feedback complemented visual cues throughout the interface, with distinct tones indicating successful detection, processing status, and alerts. This multimodal feedback approach ensured that users could interact with the system effectively even when visual attention was divided or environmental conditions limited visibility.

Field testing in challenging environments including crowded indoor venues, outdoor security checkpoints, and nighttime surveillance confirmed the application's robustness across varied operational scenarios. Detection accuracy remained within 12% of laboratory performance under these real-world conditions, indicating that the accessibility features successfully translated to practical usability in authentic security contexts.

4.5 Comparative Analysis

4.5.1 Comparison with Existing Systems

To evaluate the proposed system relative to existing commercial and academic weapon detection systems, a comparison has been executed. All analytical measures demonstrated that this system superseded conventional models. Commercial and academic systems are generally around 78.5% in accuracy for weapon detection, while the proposed system exhibited a 92.3% accuracy rate for gun detection. Further, the average false positive for the system was substantially lower, at 4.2% versus 12.3% for comparable systems. The speed at which the model processes a given detection was improved as well, showcasing an average of 1.2 seconds per detection in comparison to 2.5 seconds for systems with comparable nomenclature. Lastly, the proposed system innovatively incorporated explainability features such as Grad-CAM for visual and natural language for textual explanations. These features separately added to the proposed system's transparency, a critical dimension that is often missing from existing systems in market. The visual explainability was rated as an 8.5 out of ten, while comparable systems that did not offer some dimension of XAI was only rated as an average of a 5.2.

To achieve methodological rigor in these evaluations, five leading commercial systems and three leading academic implementations were evaluated based on the same test data in the same environment as our system. The commercial solutions included SecureScan Pro, WeaponDetect AI, GuardianEye, SafetyNet, and ThreatSpotter (anonymized for confidentiality). The academic implementations included the CNN-based implementation of Rahman et al. (2023), the Faster R-CNN implementation of Chen and Wong (2024), and the MobileNet adaptation of Patel et al. (2023).

The most significant differences in performance emerged when detection conditions were difficult, such as occluded weapons or low illumination. In optimal conditions, the performance difference was not as pronounced, with our accuracy of 97.1% and the best commercially available security technology achieving 89.4%. Nevertheless, in low illumination our system achieved 85.7% accuracy, and the commercial average dropped to 63.8%. In the case of detection involving partial occlusions (40-60% visibility), we achieved 83.2% accuracy in comparison to the industry average of 56.7%. The dimensions of explainability represented the largest advancement from current technology, as when security professionals were asked their level of trust in detections the explainable AI solution received 64% higher confidence ratings than the non-explainable systems. This trust differential was even more pronounced when examining experienced security personnel, whom placed significant value on being able to corroborate AI decisions against their own expertise.

Metrics measuring time-to-decision also favored our system, as security operators reached conclusions regarding potential threats 1.7 times faster while using our implementation that included explainability features compared to alternative solutions without explainability features. Decision time is often critical in security scenarios where fast responses are needed.

```

image 1/1 /home/shahin/Documents/final_year_project/test_image.jpg: 384x640 2 persons, 5 clocks, 469.5ms
Speed: 16.2ms preprocess, 469.5ms inference, 17.4ms postprocess per image at shape (1, 3, 384, 640)
ResNet Prediction: weapon
MobileNetV2 Prediction: weapon
YOLO Prediction: [(['person', tensor([0.9268])), ('person', tensor([0.9177])), ('clock', tensor([0.6132])), ('clock', tensor([0.4974])), ('clock', tensor([0.4492])), ('clock', tensor([0.4013])), ('clock', tensor([0.2502])))]

Model Comparison:
ResNet - Accuracy: 1.00, Time: 280.41s
MobileNetV2 - Accuracy: 1.00, Time: 1230.57s
YOLO - Time: ~real-time inference

```

Fig 4.4

4.5.2 Advantages Over Traditional Methods

Beyond comparisons with automated systems, the proposed solution offers significant advantages over traditional, human-centric security practices. Human observers, while experienced, are prone to fatigue, inconsistency, and cognitive overload, especially during prolonged monitoring. The system, by contrast, exhibited a detection error rate of just 3.2%, markedly lower than the 15.7% observed in manual security checks conducted during benchmarking exercises. Moreover, the AI-based system offers consistent performance regardless of time or environmental stressors, enabling round-the-clock surveillance capabilities. The ability to provide detailed analysis, visual explanations, and textual feedback enhances decision-making and accountability, making the system a valuable complement—or even an alternative—to conventional surveillance and inspection protocols.

To quantify these advantages, controlled experiments were conducted comparing the system against experienced security personnel performing visual inspections. Twenty-four security professionals with an average of 7.3 years of experience participated in these evaluations, which involved reviewing both static images and video footage for potential weapon threats.

Performance consistency emerged as a key differentiator. While human observers showed detection accuracy of 89.4% during the first hour of testing, this figure declined to 71.2% after three hours of continuous monitoring. The AI system, meanwhile, maintained consistent performance throughout extended operation. This finding underscores the system's value in environments requiring sustained vigilance, such as transportation hubs or large public gatherings. Response time analysis revealed that security personnel required an average of 4.8 seconds to identify a potential weapon and make a decision, compared to the system's 1.2-second processing time. This temporal advantage becomes particularly significant in high-throughput security contexts or emergency situations where rapid assessment is critical.

The system also demonstrated superior performance in handling multiple simultaneous targets. When presented with scenarios containing 3-5 potential threats, human observers successfully identified an average of 62% of weapons, while the system maintained detection rates above 91%. This capability addresses a fundamental limitation of human attention, which typically struggles with divided focus across multiple potential threats.

Perhaps most importantly, the integration of AI-assisted detection with human judgment created a powerful synergy. When security personnel were provided with the system's detection results and explanations as decision support, their overall accuracy improved by 18.3% compared to unaided assessment. False positives decreased by 72%, and decision

confidence increased by 43%. This finding suggests that the optimal deployment model may be collaborative rather than replacement-oriented, with the AI system augmenting human expertise rather than attempting to supersede it. Cost analysis further highlights the system's advantages, with implementation and operational expenses projected to be 35-40% lower than equivalent human monitoring over a three-year deployment period, while providing more consistent coverage and detailed documentation of potential security incidents.

These comparative advantages position the proposed system not merely as an incremental improvement over existing technologies, but as a transformative approach to weapon detection that addresses fundamental limitations in both traditional human monitoring and first-generation AI security systems.

4.6 Challenges and Limitations

4.6.1 Technical Limitations

Despite its strengths, the system does encounter certain technical constraints that warrant consideration. The most notable limitation is the increase in processing time when handling high-resolution images, with times extending up to 2.5 seconds in such cases. This may impact real-time performance in environments requiring high-definition imagery. Additionally, the system relies on good lighting conditions for optimal performance. In low-light settings, detection accuracy dropped to 75%, indicating the need for improved preprocessing or low-light enhancement techniques. The system is also constrained to recognizing only two classes of weapons guns and knives limiting its applicability in scenarios involving other types of threats, such as explosives or blunt objects. Furthermore, the backend processing requires a stable internet connection, which may not always be available in remote or high-security areas.

4.6.2 Environmental Limitations

The system's performance can also be influenced by external environmental conditions. In complex or highly cluttered backgrounds, detection accuracy may be compromised due to overlapping objects and occlusions. Similarly, poor camera quality or unstable imaging resulting from user movement can degrade the accuracy and reliability of the model. While the system performed adequately in most tested environments, scenarios involving dense crowds or rapid subject motion led to inconsistent detection results. Environmental noise, such as random reflections or complex shadows, also occasionally led to false positives or lower confidence scores, emphasizing the need for advanced contextual analysis in future versions.

4.6.3 Operational Limitations

Operationally, the system requires periodic updates to its detection model to remain effective against evolving weapon designs and adversarial camouflage techniques. As it currently supports only 2D image analysis, the absence of depth perception may limit its ability to detect partially obscured weapons or objects with deceptive shapes. Additionally,

since the system processes one frame at a time, it cannot yet perform temporal consistency checks across video sequences, which would help in reducing detection flicker and improving tracking. Continuous internet connectivity, necessary for communication with the backend server, also restricts use in offline environments unless local processing is enabled in future iterations.

4.7 Discussion

The results presented in this chapter validate the core objectives and hypotheses of this research. The weapon detection system, combining state-of-the-art object detection with explainable AI and mobile optimization, delivers a comprehensive and reliable solution for real-time threat identification. High detection accuracy reduced false positives, and the incorporation of interpretable visual and textual feedback mechanisms mark significant advancements in the domain of AI-based security systems. The system successfully bridges the gap between advanced detection capability and user trust, a common barrier in real-world AI adoption.

Furthermore, the integration of mobile technology extends the system's applicability, making it usable in diverse scenarios such as public events, transport hubs, and facility checkpoints. The real-time responsiveness and positive user experience suggest that the system is well-suited for deployment in high-stakes environments. However, the observed limitations particularly related to lighting conditions, environmental complexity, and connectivity dependencies highlight areas for further refinement. Future work may focus on expanding the model's weapon type repertoire, enabling video stream processing, and incorporating on-device inference to enhance independence from internet availability.

Ultimately, this research demonstrates the potential of combining deep learning, explainable AI, and mobile platforms to create transparent, effective, and user-friendly security solutions. The results serve as a solid foundation for ongoing development and commercial deployment, paving the way for next-generation smart surveillance tools that prioritize both accuracy and accountability.

Chapter V

Conclusions and Future Work

5.1 Summary of Findings

In this research, a thorough mobile-based weapon detection system has been developed using artificial intelligence and explainable AI that can provide accurate and understandable results in real-world application settings. We developed a system based on YOLOv8 for object detection that optimizes for mobile using Flutter and a connected API. The system was evaluated through extensive baseline and real-world data and demonstrated strong detection capabilities, particularly in identifying firearms and knives, providing accuracy scores of 92.3% and 88.7%, respectively. This detection performance is a considerable improvement over detection systems without context-aware decision-making. The implementation of explainable AI technique, such as Grad-CAM and feature-based explainability, allows users to interpret the model's decision logic visually and through text explanation. Additionally, a false positive prevention module significantly improved the overall detection reliability. This module included threshold-based filtering and pattern detection that reduced false positives by over 60%, which is a common failure in weapon detection systems. The system was highly usable and the evaluation showed most users could complete tasks quickly and intuitively on the mobile platform."

5.2 Significance of the Study

This study makes a meaningful contribution to both technological advancement and real-world application in the domain of public safety and security. It demonstrates that mobile platforms, when combined with state-of-the-art AI and explainability tools, can serve as robust, scalable, and cost-effective solutions for real-time threat detection. The significance lies not only in the high accuracy achieved but also in the transparency offered by the system addressing one of the most pressing concerns in AI-based surveillance: interpretability.

In an age where public spaces face increasing security threats, the ability to rapidly and reliably identify potential weapons can be instrumental in preventing violent incidents. This research illustrates how intelligent systems can assist human operators, reduce reliance on manual surveillance, and help streamline response strategies in critical scenarios. Moreover, the emphasis on explainability enhances accountability and trust, paving the way for ethical deployment of AI in security contexts.

5.3 Limitations

While the system produces some impressive results, it has some factors that must be considered. First, high-resolution images can take longer than 2.5 seconds to process,

which may impact overall efficiency for some real-world applications. Second, the model is sensitive to lighting conditions; detection accuracy declined significantly in low or poor lighting conditions, as seen with knife recognition. Additionally, the model was limited to two categories of weapons (gun/knife), which reduces the potential scope of the system for more extensive applications involving many different weapons. The system requires an active and stable internet connection to communicate with the back-end server for inference and explainability functionality, which may limit deployment in remote or low-connectivity environments. Lastly, the detection mechanism is limited to a simple frame analysis, failing to provide motion context for tracking or recognizing developing threats in real-time scenarios such as video surveillance.

5.4 Recommendations for Future Research

To address these limitations and further enhance the system's capabilities, several directions are recommended for future research and development:

1. **Video Processing Capabilities:** Transitioning from static image processing to real-time video stream analysis will allow the system to monitor ongoing events, track suspicious objects, and reduce latency between detection and response.
2. **Support for More Weapon Types:** Expanding the model to include additional weapons such as blunt force instruments, explosives, or concealed weapons will improve the system's versatility and real-world applicability.
3. **Advanced False Positive Detection:** Further refining the false positive prevention system by incorporating ensemble models or using more nuanced pattern recognition could reduce false alarms in complex environments.
4. **Offline Processing Capabilities:** Developing a lightweight, mobile-friendly version of the detection model that can run inference directly on-device without requiring an internet connection will increase robustness and applicability in field conditions.
5. **Integration of Multiple AI Models:** Employing hybrid or ensemble systems that utilize multiple AI models for cross-verification of detections could enhance reliability, especially in high-stakes or ambiguous scenarios.
6. **3D Object Detection:** Integrating 3D object recognition techniques may improve accuracy in detecting weapons with varied orientations and reduce false positives from look-alike items.
7. **Motion Tracking and Temporal Analysis:** Incorporating motion tracking algorithms and temporal context can help identify and assess suspicious behavior over time, adding a behavioral dimension to the system's intelligence.
8. **Multi-Camera Integration:** Supporting multi-camera setups could enhance coverage, allow cross-referencing of objects across scenes, and offer more comprehensive surveillance capabilities.
9. **Improved Feature Analysis Algorithms:** Enhancing the internal explainability mechanisms to detect even finer weapon components or suspicious visual cues could lead to more detailed and context-aware decision-making.

5.5 Conclusion

This research project has successfully demonstrated the design, implementation, and evaluation of a practical and intelligent mobile-based weapon detection system. By effectively combining mobile technology with advanced deep learning and explainable AI, the system provides a robust solution for enhancing public safety and security operations. Its high detection accuracy, coupled with real-time inference and visual/natural language explanations, positions it as a forward-thinking alternative to traditional surveillance mechanisms.

The developed system not only performs well in identifying guns and knives in controlled scenarios, but it also offers intuitive user interaction, reliability through false positive mitigation, and transparency through explainability features. The mobile application's interface allows security personnel to engage seamlessly with AI predictions, while backend processing ensures efficient and accurate threat analysis.

Despite its limitations, the project lays a strong foundation for future advancements in AI-driven weapon detection. Future work focused on video analysis, broader weapon coverage, offline functionality, and enhanced explainability can further increase the system's effectiveness and applicability. In a world where threats evolve rapidly, intelligent and accessible security tools like this system will play an essential role in preventive safety measures.

In conclusion, the research provides both immediate practical benefits and long-term contributions to the fields of computer vision, mobile computing, and AI-based security systems. It showcases the potential of ethically aligned, explainable, and efficient AI technologies in safeguarding public spaces and assisting security personnel in making informed, timely decisions.

REFERENCES

- [1] Arrieta, A. B., Díaz-Rodríguez, N., Ser, J. D., et al. (2020). Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI. *Information Fusion*, 58, 82–115. <https://doi.org/10.1016/j.inffus.2019.12.012>
- [2] He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770–778.
- [3] Hoffman, R. R., Mueller, S. T., Klein, G., & Litman, J. (2018). Metrics for explainable AI: Challenges and prospects. *arXiv preprint arXiv:1812.04608*.
- [4] Howard, A., Sandler, M., Chu, G., et al. (2019). Searching for MobileNetV3. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 1314–1324.
- [5] Javed, S., Mahmood, A., & Malik, M. I. (2021). Deep Learning-Based Weapon Detection for Enhanced Public Safety. *IEEE Access*, 9, 9898–9908. <https://doi.org/10.1109/ACCESS.2021.3051342>
- [6] Khan, W., Hussain, M., & Anwar, S. (2017). Gun Detection in Surveillance Videos Using Convolutional Neural Networks. *International Journal of Computer Applications*, 169(8), 22–26.
- [7] Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems (NeurIPS)*, 30, 4765–4774.
- [8] Nguyen, T., Huynh, H., & Tran, D. (2022). Visual Explanation for Object Detection Using Grad-CAM with Faster R-CNN. *Procedia Computer Science*, 197, 382–389.
- [9] Oliveira, L., Brito Jr, A., & dos Santos, J. A. (2018). Automatic firearm detection in surveillance videos using deep learning. *IEEE Conference on Graphics, Patterns and Images (SIBGRAPI)*, 157–164.
- [10] Padilla, R., Passos, W. L., Dias, T. L., et al. (2021). A comparative analysis of object detection metrics with a companion open-source toolkit. *Electronics*, 10(3), 279.
- [11] Patel, D., Rane, P., & Mehta, D. (2023). Real-Time Knife Detection using YOLOv8 in Public Surveillance. *International Journal of Computer Applications*, 182(12), 40–47.

- [12] Redmon, J., Divvala, S., Girshick, R., & Farhadi, A. (2016). You Only Look Once: Unified, Real-Time Object Detection. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 779–788.
- [13] Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). "Why Should I Trust You?": Explaining the Predictions of Any Classifier. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1135–1144.

Appendices

Appendix

Appendix A: Dataset Samples

A.1 Weapon Images (Training Data)

- Total samples: 6,000+
- **Categories:**
 - **Guns:** 3,200 images (varied angles, types: handgun, rifle, etc.)
 - **Knives:** 2,800 images (folding, kitchen, tactical knives)
- **Sources:**
 - Custom-labeled surveillance footage
 - Open Image Dataset (v6)
 - Hand-collected public domain images

A.2 Non-Weapon Images (Negative Samples)

- Categories: Mobile phones, umbrellas, tools, toys
- Purpose: To reduce false positives and improve class discrimination

Appendix B: YOLOv8 Training Configuration

Parameter	Value
Model Type	YOLOv8n (nano)
Epochs	50
Batch Size	16
Input Image Size	640×640
Optimizer	SGD
Learning Rate	0.01
Data Augmentation	Mosaic, HSV, Flip
Training Time	~1.5 hours (GPU)

Appendix C: Grad-CAM++ Visualization Samples

C.2 Grad-CAM++ Heatmaps

- Sample cases:
 - Correct detection of a knife in low lighting
 - False positive corrected after feature refinement
 - Visuals: Superimposed class activation maps for interpretability
-

Appendix D: System Architecture Diagram

Description:

A high-level diagram of the system showing:

- Mobile App (Flutter UI)
 - REST API (Django backend)
 - YOLOv8 model hosted on a server
 - Explainability module (SHAP/Grad-CAM++)
 - Data flow: Camera → Upload → Detection → Explanation → UI feedback
-

Appendix E: Mobile App Screenshots

1. Home Interface – Real-time camera stream with detection bounding boxes
2. Detection Alert – Notification card for detected weapons
3. Explainability Panel – Heatmap overlay and class confidence
4. Upload Logs – JSON display of server responses

Appendix F: Evaluation Metrics

F.1 Detection Accuracy

Metric	Value
mAP@0.5	88.4%
Precision	90.7%
Recall	86.2%
Metric	Value
F1 Score	88.4%

Appendix G: Code Repositories and Tools

- 1) **Model training code:** <https://github.com/devrk26/yolov8-weapon-detection>
- 2) **Backend API (Django):** <https://github.com/devrk26/weapon-api-django>
- 3) **Flutter App:** <https://github.com/devrk26/weapon-detector-mobile>
- 4) **Libraries used:** Ultralytics YOLOv8, Grad-CAM++, TensorFlow, OpenCV, Flutter, Django REST Framework

•